



# PEP 603 – Adding a frozenmap type to collections

**Author** Yury Selivanov <yury at edgedb.com>

**Status** Draft

**Type** Standards Track

**Created** 12-Sep-2019

**Python-Version** 3.9

**Post-History** 12-Sep-2019

## ► Table of Contents

## Abstract

A *persistent data structure* is defined as a data structure that preserves the previous version of the data when the data is modified. Such data structures are effectively *immutable*, as operations on them do not update the structure in-place, but instead always yield a new updated structure (see [0] for more details.)

This PEP proposes to add a new fully persistent and immutable mapping type called `frozenmap` to the `collections` module.

The bulk of `frozenmap`'s reference implementation is already used in CPython to implement the `contextvars` module.

## Rationale

Python has two immutable collection types: `tuple` and `frozenset`. These types can be used to represent immutable lists and sets. However, a way to represent immutable *mappings* does not yet exist, and this PEP proposes a `frozenmap` to implement an immutable *mapping*.

The proposed `frozenmap` type:

- implements the `collections.abc.Mapping` protocol,
- supports pickling, and
- provides an API for efficient creation of “modified” versions.

The following use cases illustrate why an immutable mapping is desirable:

- Immutable mappings are hashable which allows their use as dictionary keys or set elements.

This hashable property permits functions decorated with `@functools.lru_cache()` to accept immutable mappings as arguments. Unlike an immutable mapping, passing a plain `dict` to such a function results in error.

- Immutable mappings can hold complex state. Since immutable mappings can be copied by reference, transactional mutation of state can be efficiently implemented.
- Immutable mappings can be used to safely share dictionaries across thread and asynchronous task boundaries. The immutability makes it easier to reason about threads and asynchronous tasks.

Lastly, CPython [1] already contains the main portion of the C code required for the `frozenmap` implementation. The C code already exists to implement the `contextvars` module (see PEP 567 for more details.) Exposing this C code via a public collection type drastically increases the number of users of the code. This leads to increased code quality by discovering bugs and improving performance which without a `frozenmap` collection would be very challenging because most programs use the `contextvars` module indirectly.

## Specification

A new public immutable type `frozenmap` is added to the `collections` module.

## Construction

`frozenmap` implements a `dict`-like construction API:

- `frozenmap()` creates a new empty immutable mapping;
- `frozenmap(**kwargs)` creates a mapping from `**kwargs`, e.g. `frozenmap(x=10, y=0, z=-1)`
- `frozenmap(collection)` creates a mapping from the passed `collection` object. The passed `collection` object can be:
  - a `dict`,
  - another `frozenmap`,
  - an object with an `items()` method that is expected to return a series of key/value tuples, or
  - an iterable of key/value tuples.

## Data Access

`frozenmap` implements the `collection.abc.Mapping` protocol. Therefore, getters, membership checks, and iteration work the same way that they would for a `dict`:

```
m = frozenmap(foo='bar')

assert m['foo'] == 'bar'
assert m.get('foo') == 'bar'
assert 'foo' in m

assert 'baz' not in m
assert m.get('baz', 'missing') == 'missing'

assert m == m
assert m != frozenmap()  # m is not equal to an empty frozenmap

assert len(m) == 1

# etc.
```

## Mutation

`frozenmap` instances are immutable. That said, it is possible to efficiently produce mutated *copies* of the immutable instance.

The complexity of mutation operations is  $O(\log N)$  and the resulting `frozenmap` copies often consume very little additional memory due to the use of structural sharing (read [6] for more details.)

## `frozenmap.including(key, value)`

The method creates a new `frozenmap` copy with a new *key* / *value* pair:

```
m = frozenmap(foo=1)
m2 = m.including('bar', 100)

print(m)    # will print frozenmap({'foo': 1})
print(m2)   # will print frozenmap({'foo': 1, 'bar': 100})
```

## `frozenmap.excluding(key)`

The method produces a copy of the `frozenmap` which does not include a deleted *key*:

```
m = frozenmap(foo=1, bar=100)

m2 = m.excluding('foo')

print(m)    # will print frozenmap({'foo': 1, 'bar': 100})
print(m2)   # will print frozenmap({'bar': 1})

m3 = m.excluding('spam') # will throw a KeyError('spam')
```

## `frozenmap.union(mapping=None, **kw)`

The method produces a copy of the `frozenmap` and adds or modifies multiple key/values for the created copy. The signature of the method matches the signature of the `frozenmap` constructor:

```
m = frozenmap(foo=1)

m2 = m.union({'spam': 'ham'})
print(m2) # will print frozenmap({'foo': 1, 'spam': 'ham'})

m3 = m.union(foo=100, y=2)
print(m3) # will print frozenmap({'foo': 100, 'y': 2})

print(m)  # will print frozenmap({'foo': 1})
```

Calling the `union()` method to add/replace *N* keys is more efficient than calling the `including()` method *N* times.

## `frozenmap.mutating()`

The method allows efficient copying of a `frozenmap` instance with multiple modifications applied. This method is especially useful when the frozenmap in question contains thousands of key/value pairs and there's a need to update many of them in a performance-critical section of the code.

The `frozenmap.mutating()` method returns a mutable dict-like copy of the `frozenmap` object: an instance of `collections.FrozenMapCopy`.

The `FrozenMapCopy` objects:

- are copy-on-write views of the data of `frozenmap` instances they were created from;
- are mutable, although any mutations on them do not affect the `frozenmap` instances they were created from;
- can be passed to the `frozenmap` constructor; creating a `frozenmap` from a `FrozenMapCopy` object is an  $O(1)$  operation;
- have  $O(\log N)$  complexity for get/set operations; creating them is an  $O(1)$  operation;
- have a `FrozenMapCopy.close()` method that prevents any further access/mutation of the data;
- can be used as a context manager.

The below example illustrates how `mutating()` can be used with a context manager:

```
numbers = frozenmap((i, i ** 2) for i in range(1_000_000))

with numbers.mutating() as copy:
    for i in numbers:
        if not (numbers[i] % 997):
            del copy[i]

numbers_without_997_multiples = frozenmap(copy)

# at this point, *numbers* still has 1_000_000 key/values, and
# *numbers_without_997_multiples* is a copy of *numbers* without
# values that are multiples of 997.

for i in numbers:
    if not (numbers[i] % 593):
        del copy[i]

numbers_without_593_multiples = frozenmap(copy)

print(copy[10]) # will print 100.

print(copy[10]) # This will throw a ValueError as *copy*
                # has been closed when the "with" block
                # was executed.
```

## Iteration

As `frozenmap` implements the standard `collections.abc.Mapping` protocol, so all expected methods of iteration are supported:

```
assert list(m) == ['foo']
assert list(m.items()) == [('foo', 'bar')]
assert list(m.keys()) == ['foo']
assert list(m.values()) == ['bar']
```

Iteration in `frozenmap`, unlike in `dict`, does not preserve the insertion order.

## Hashing

`frozenmap` instances can be hashable just like `tuple` objects:

```
hash(frozenmap(foo='bar'))  # works
hash(frozenmap(foo=[]))     # will throw an error
```

## Typing

It is possible to use the standard typing notation for frozenmaps:

```
m: frozenmap[str, int] = frozenmap()
```

## Implementation

The proposed `frozenmap` immutable type uses a Hash Array Mapped Trie (HAMT) data structure. Functional programming languages, like Clojure, use HAMT to efficiently implement immutable hash tables, vectors, and sets.

### HAMT

The key design contract of HAMT is the guarantee of a predictable *value* when given the hash of a *key*. For a pair of *key* and *value*, the hash of the *key* can be used to determine the location of *value* in the hash map tree.

Immutable mappings implemented with HAMT have  $O(\log N)$  performance for `set()` and `get()` operations. This efficiency is possible because mutation operations only affect one branch of the tree, making it possible to reuse non-mutated branches, and, therefore, avoiding copying of unmodified data.

Read more about HAMT in [5]. The CPython implementation [1] has a fairly detailed description of the algorithm as well.

## Performance

## dict.copy() vs frozenmap

(10 get-item calls, 1 set-item call) x 1M iterations

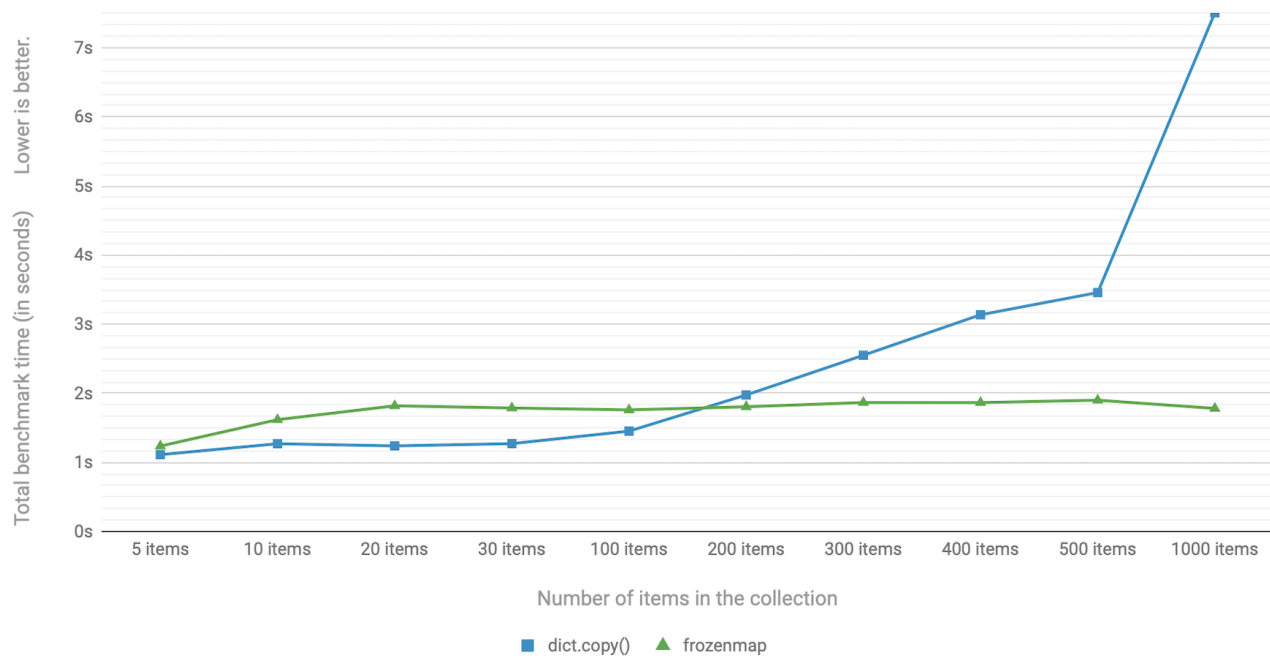


Figure 1. Benchmark code can be found here: [3].

The above chart demonstrates that:

- `frozenmap` implemented with HAMT displays near  $O(1)$  performance for all benchmarked dictionary sizes.
- `dict.copy()` becomes less efficient when using around 100-200 items.

## Lookup time: HAMT vs dict

(10 get-item calls) x 1M iterations

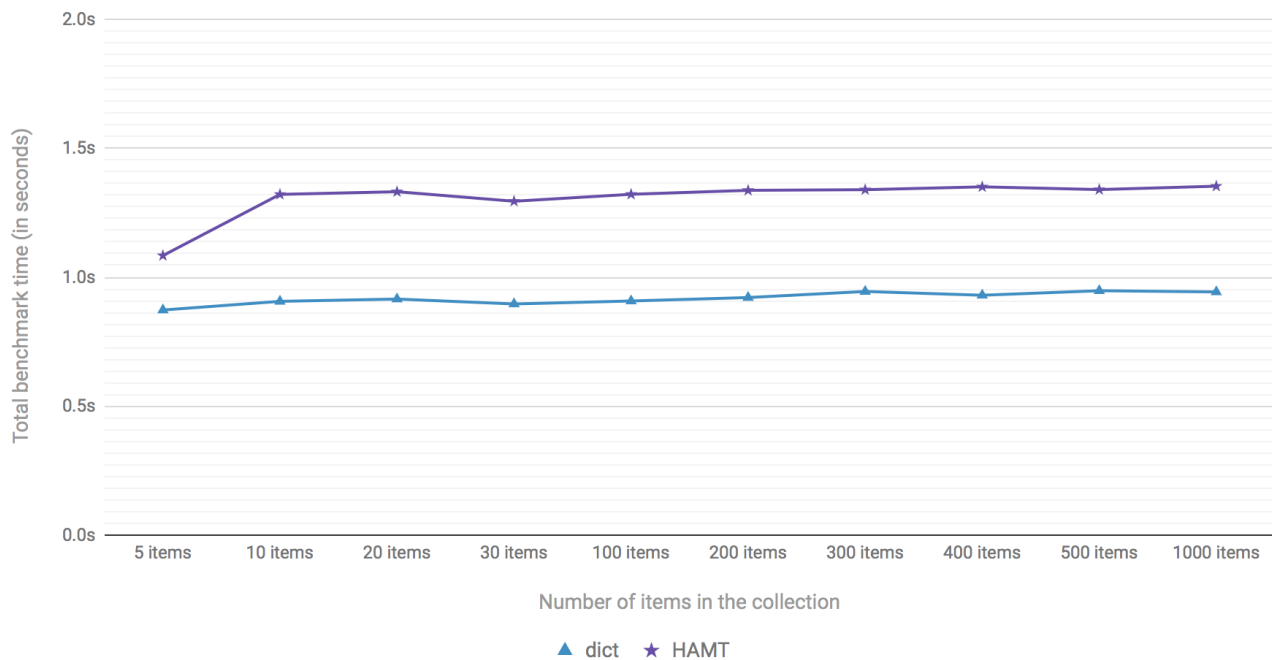


Figure 2. Benchmark code can be found here: [4].

Figure 2 compares the lookup costs of `dict` versus a HAMT-based immutable mapping. HAMT lookup time is ~30% slower than Python `dict` lookups on average. This performance difference exists since traversing a shallow tree is less efficient than lookup in a flat continuous array.

Further to that, quoting [6]: “[using HAMT] means that in practice while insertions, deletions, and lookups into a persistent hash array mapped trie have a computational complexity of  $O(\log n)$ , for most applications they are effectively constant time, as it would require an extremely large number of entries to make any operation take more than a dozen steps.”

## Design Considerations

### Why “frozenmap” and not “FrozenMap”

The lower-case “frozenmap” resonates well with the `frozenset` built-in as well as with types like `collections.defaultdict`.

### Why “frozenmap” and not “frozendict”

“Dict” has a very specific meaning in Python:

- a dict is a concrete implementation of `abc.MutableMapping` with  $O(1)$  get and set operations (`frozenmap` has  $O(\log N)$  complexity);
- Python dicts preserve insertion order.

The proposed `frozenmap` does not have these mentioned properties. Instead, `frozenmap` has an  $O(\log N)$  cost of set/get operations, and it only implements the `abc.Mapping` protocol.

# Implementation

The full implementation of the proposed `frozenmap` type is available at [2]. The package includes C and pure Python implementations of the type.

See also the HAMT collection implementation as part of the CPython project tree here: [1].

## References

[0]	<a href="https://en.wikipedia.org/wiki/Persistent_data_structure">https://en.wikipedia.org/wiki/Persistent_data_structure</a>
[1] (1, 2, 3)	<a href="https://github.com/python/cpython/blob/3.8/Python/hamt.c">https://github.com/python/cpython/blob/3.8/Python/hamt.c</a>
[2]	<a href="https://github.com/MagicStack/immutables">https://github.com/MagicStack/immutables</a>
[3]	<a href="https://gist.github.com/1st1/be5a1c10aceb0775d0406e879cf87344">https://gist.github.com/1st1/be5a1c10aceb0775d0406e879cf87344</a>
[4]	<a href="https://gist.github.com/1st1/dbe27f2e14c30cce6f0b5fddfc8c437e">https://gist.github.com/1st1/dbe27f2e14c30cce6f0b5fddfc8c437e</a>
[5]	<a href="https://en.wikipedia.org/wiki/Hash_array_mapped_trie#cite_note-bagwell-1">https://en.wikipedia.org/wiki/Hash_array_mapped_trie#cite_note-bagwell-1</a>
[6] (1, 2)	<a href="https://en.wikipedia.org/wiki/Persistent_data_structure#Trees">https://en.wikipedia.org/wiki/Persistent_data_structure#Trees</a>

## Acknowledgments

I thank Carol Willing, Łukasz Langa, Larry Hastings, and Guido van Rossum for their feedback, ideas, edits, and discussions around this PEP.

## Copyright

This document is placed in the public domain or under the CC0-1.0-Universal license, whichever is more permissive.

Source: <https://github.com/python/peps/blob/main/pep-0603.rst>

Last modified: 2022-03-11 20:41:57 GMT