



CSC 139

Operating System Principles

Chapter 3

Process

Herbert G. Mayer, CSU CSC

8/1/2022

Copied with permission from: Silberschatz, Galvin and Gagne © 2013

Syllabus

- **Process** Concept
- **Process Scheduling**
- **Process States**
- **Operations on Processes**
- **Interprocess** Communication (IPC)
- **Examples of IPC Systems**
- **Client-Server Systems**

Objectives


- Introduce **process**: Process is a **program in execution**, which forms the basis for computations, managed by an OS
- Describe various needs and features of processes: **Scheduling**, priorities, creation, execution, interruption, termination, and communication
- Explore **interprocess communication** using **shared memory** or **message passing**
- Sketch out communication in **client-server** systems
- Discuss industry- and academic standards that are process related!

Process Concept

Informally:

- Process is sometimes mistaken to be **synonymous to program**; but they do differ
- **Program & process** do share a key ingredient: **object code**
- Object code per-se is **lifeless**, just a collection of bits; some bits are data, others are executable code: but the key point is: **executable**, not yet **executing**! Not loaded!
- **Process** is a program that is **executing**, code that is currently running, has been loaded and is **ready to run** but is in the wait queue!
- To actually execute, a process needs resources e.g.: **processor**, loaded **object** code, **memory** to run in, may need **files** for reading data and generating data
- All managed under OS control!

Windows Process Manager



Process Concept

More Formally:

- **Process** is a **program in execution** under OS control
- **OS** executes and manages processes in various modes:
 - **Supervisor mode** – to manage OS's resources, user processes
 - **Batch mode** – user jobs to be executed w/o user interaction
 - **Time-shared mode** – user interacts with running software
- Various **process** components:
 1. Program code, AKA **text** section (in antique ☺ IBM terms)
 2. HW resources, including **registers**, static **memory**, **stack**, **heap**
 3. **Stack** containing temporary & local data, varies in size during run, grows at call, shrinks at return, starts and ends at ~0 size
 4. Functions & **function parameters**, return address, local objects
 5. **Data** section containing global objects; possibly shared
 6. Files for data to be read (input) and generated (output)

Process Concept


Remember 😊

Process is a program in execution!

- A **program** *per-se* is not alive, it is just a sequence of silent bits, stored somewhere, perhaps on disk (or solely in a programmer's mind 😊)
- Once a program is loaded into memory, ready to run, or actually **running on a system**, then it is **alive**, an active process, consuming resources, inhabiting a position in some scheduling queue, lasting a certain time, consuming input, producing output
- Then the program has become a **process**
- A **process is alive**
- A program only has the potential to live in the future; even a process in wait state is **alive**

Process Concept


Batch Process, involves human operator; *in days of old*



Process vs. Program

- **Program** is a *passive* entity, just bits (presumably meaningful) stored on a disk, as an executable file
- **Process** is *active, is alive, is an executing program*
 - Program becomes process when executable file is loaded into memory, *ready to be* granted processor time
- Execution of program starts via 1. GUI mouse click, or 2. Command line entry of its name, or 3. Pushing a button to read cards, or 4. Touching an icon, etc.
- One program can evolve into several processes in various ways:
 - Consider multiple users executing the same program; such as an editor on a time-shared main frame: distinct processes
 - Or single user program, when executing, may fork into multiple sub-processes; desirable on MP target for speed

Process in Memory

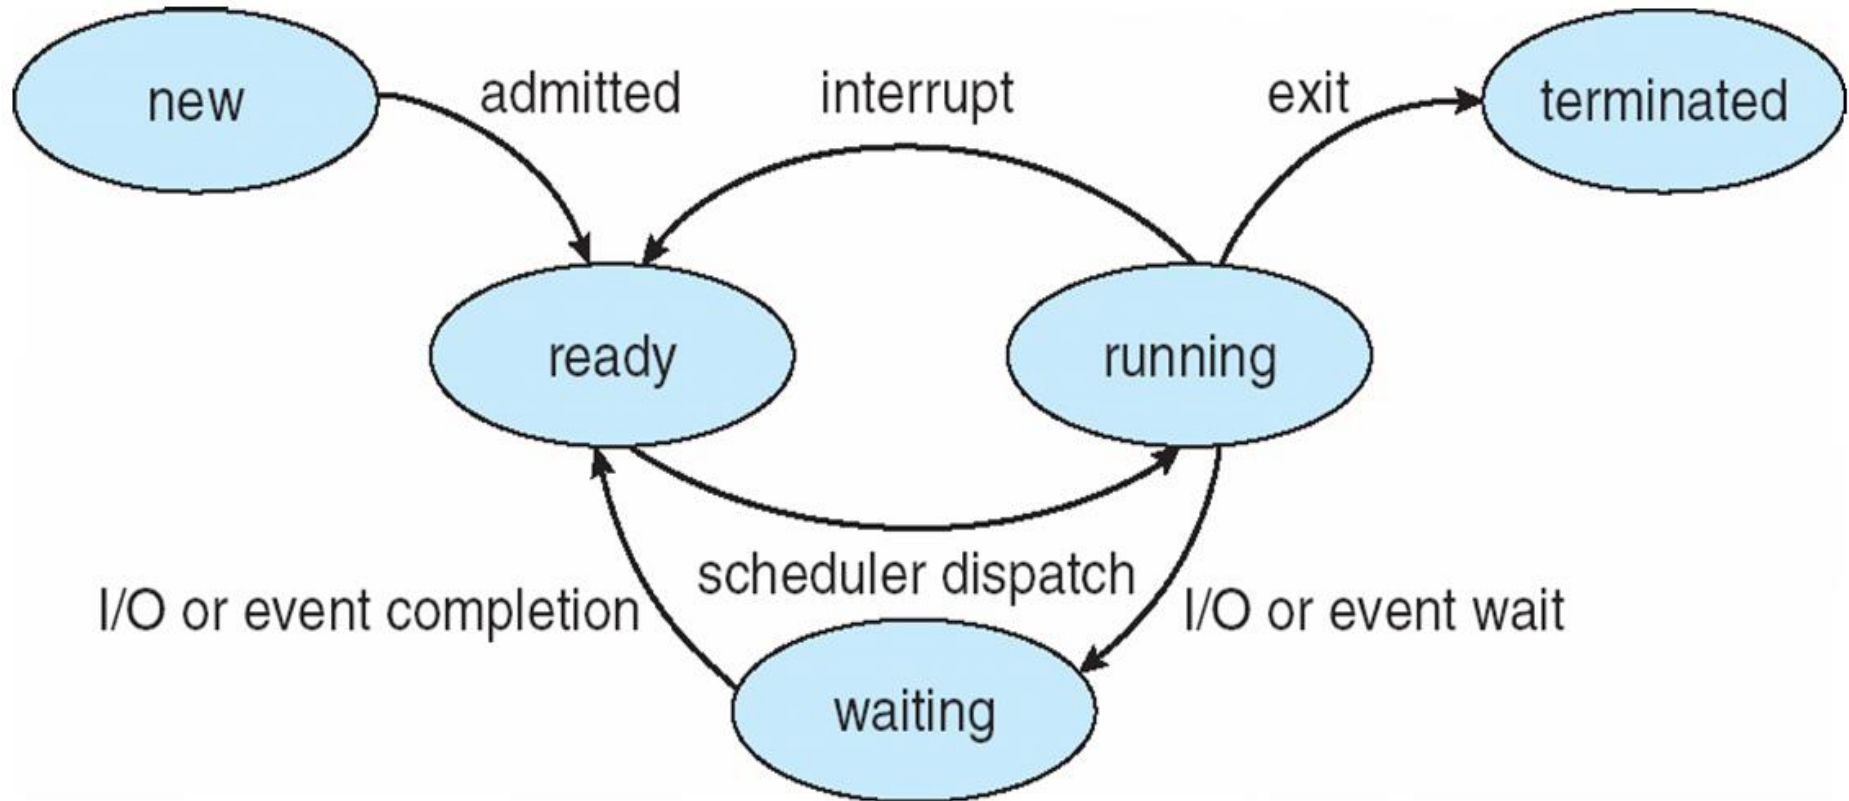


Process States

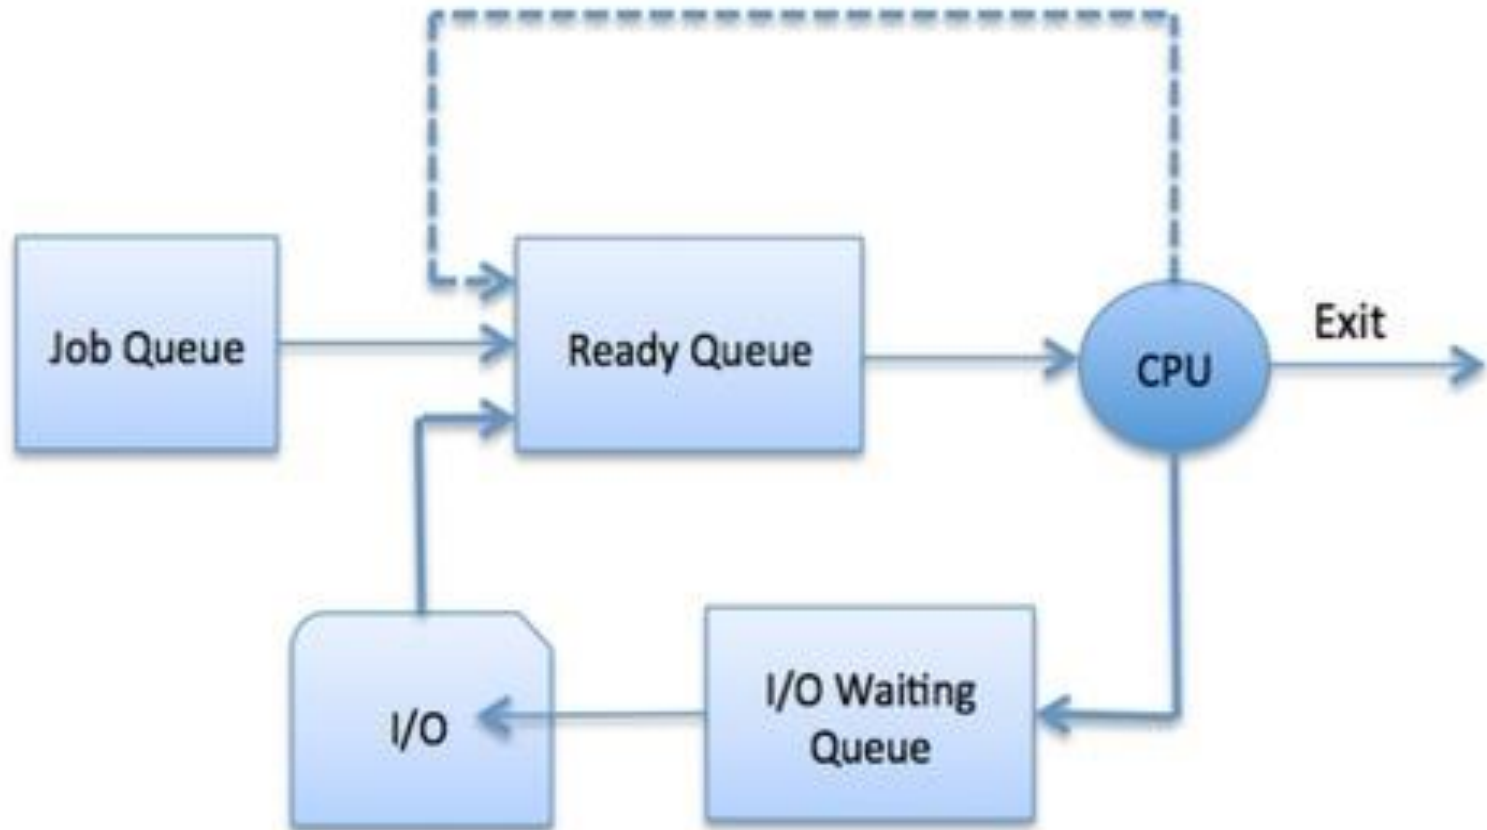
While a process executes, over time it changes **state**

- **new**: Process is created by loading from disk into memory
- **running**: Instructions are being executed on a CPU
- **waiting**: Process is waiting for **some event to occur** on its behalf; could be data input, or another event; may change priority while waiting
- **ready**: Process is waiting to be assigned to a processor, i.e. it is ready run but not yet running, as it has not been granted a CPU
- **terminated**: Process has completed execution; its whole memory space is being freed; all other resources will be re-acquired by OS; PID will be purged
- A priori and a posteriori: the **OS owns it all!**

Process States

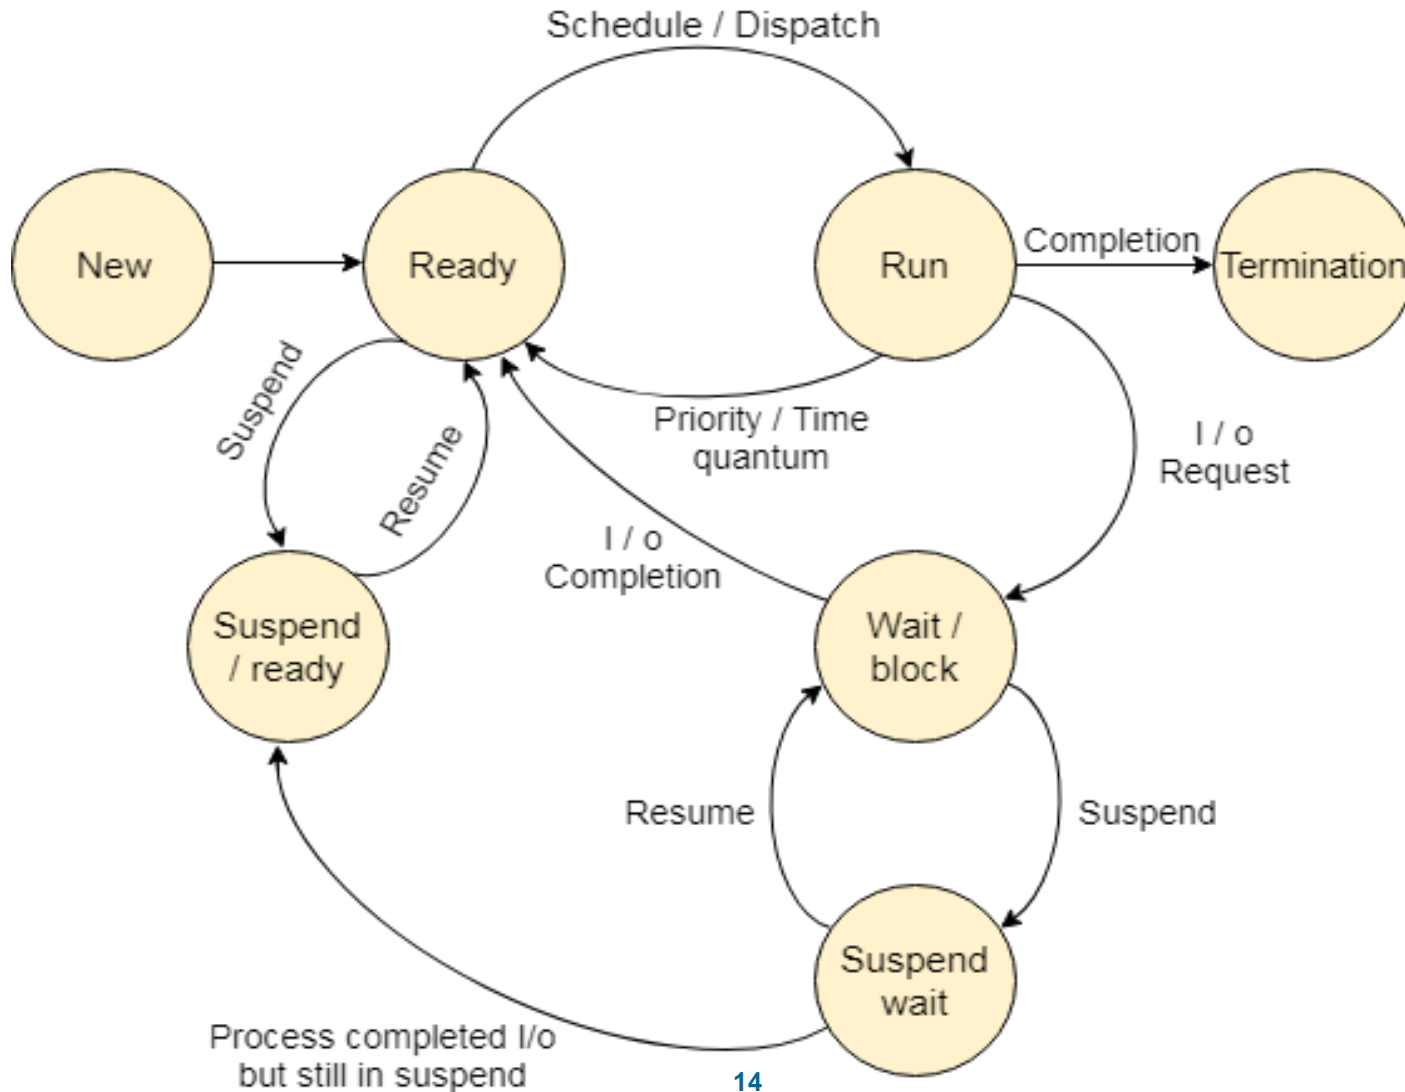


Process Queue



Process States

Same Process State Model, More Detail

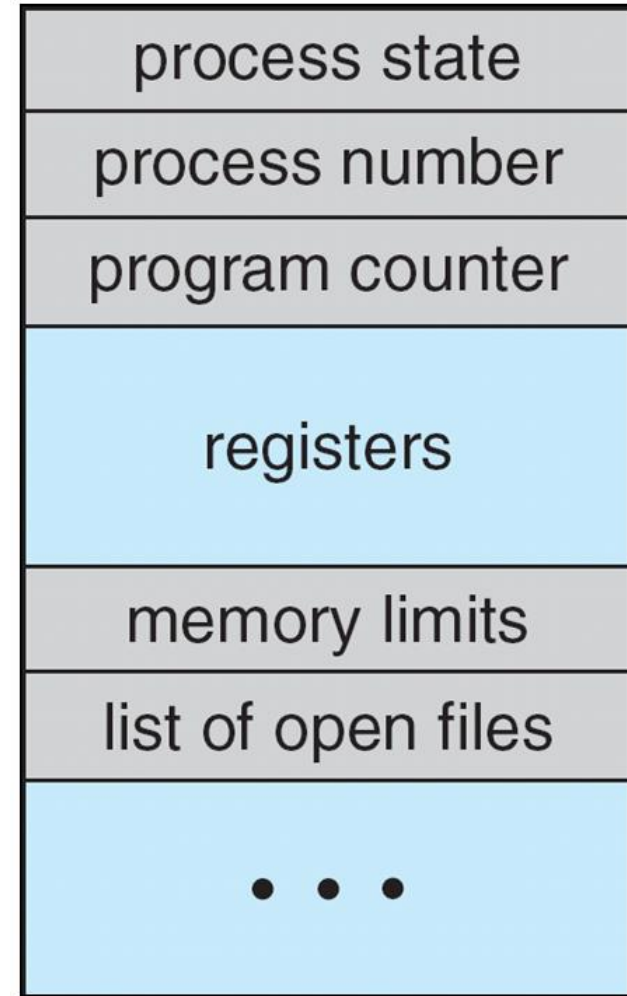


Process Control Block (PCB)

Information associated with process:
e.g. an actual C++ data structure:

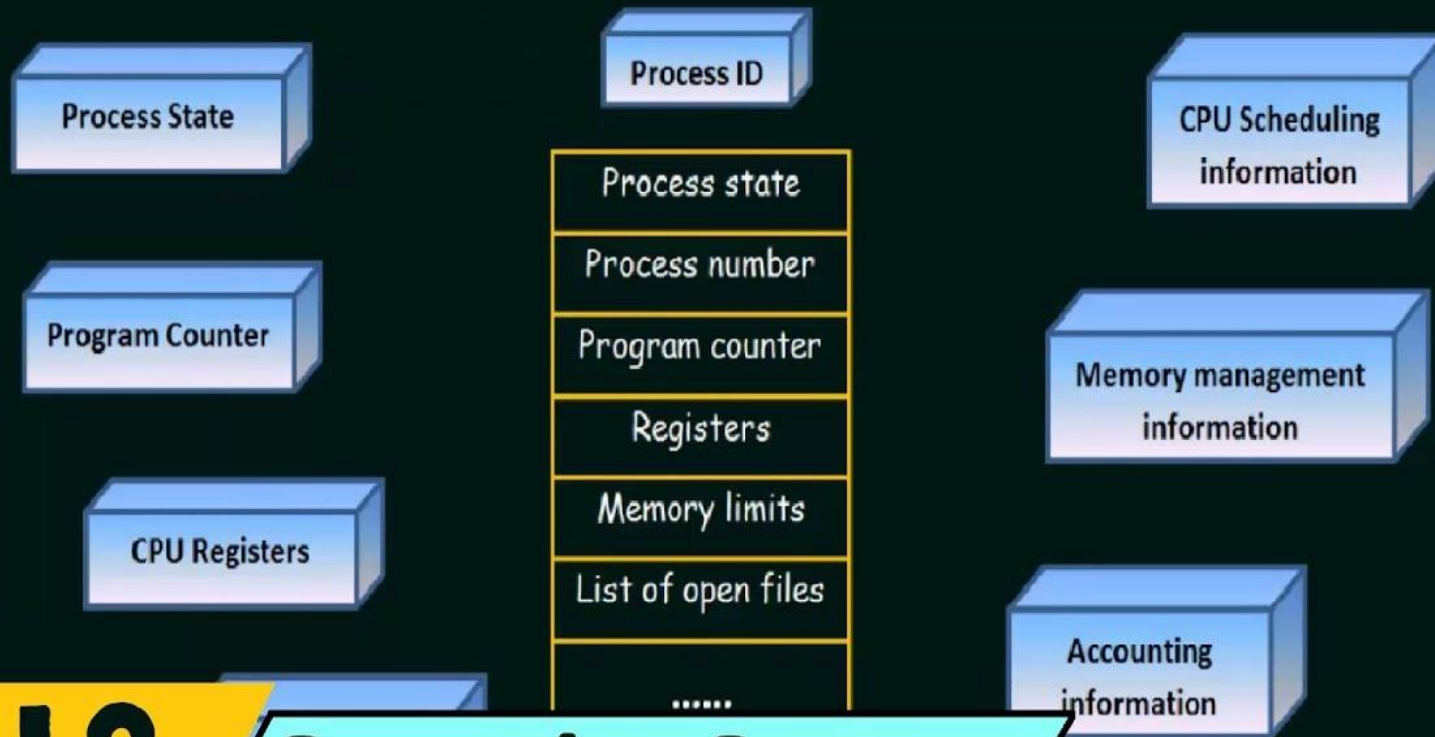
AKA **process control block** PCB

- Process ID (AKA **PID**): unique number
- Process states: **run**, **wait** or **suspend**
- Program counter – address of next instruction to execute (AKA **pc** or **ip**)
- CPU **registers** – contents of all process-centric registers; generally: **all regs**
- CPU scheduling information – priority, scheduling queue pointers
- Memory management information – memory allocated to process
- Accounting info – **CPU use**, clock time elapsed since start; time limits if defined
- I/O status information – I/O devices allocated to process, list of **open files**



PCB

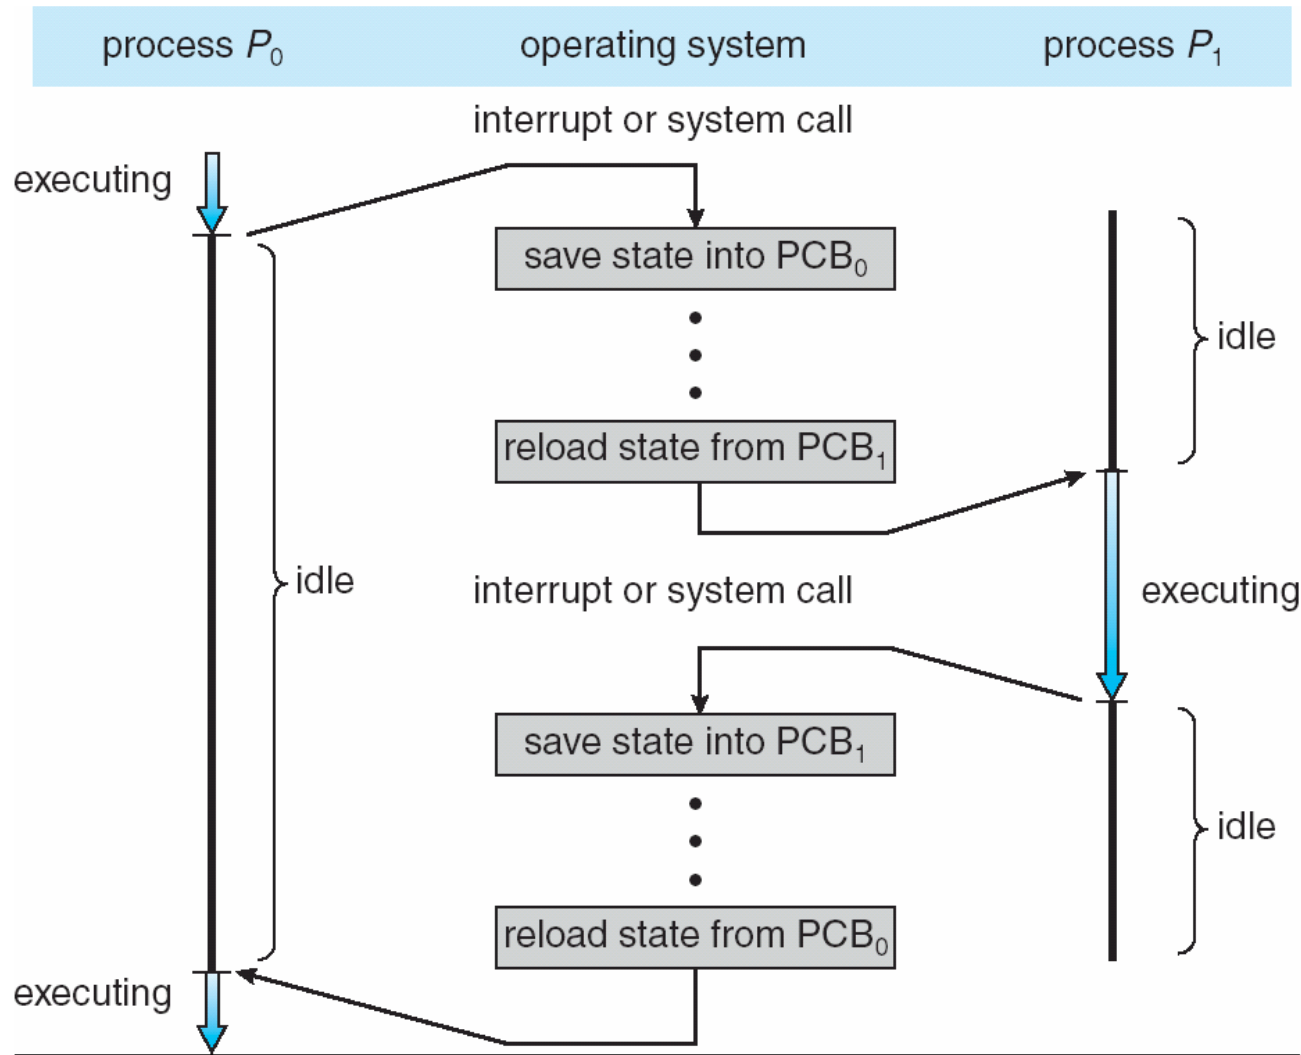
Process Control Block



18

Operating System

CPU Switch *Process to Process*



CPU Switch *Process to Process*

- Visible from above picture that **inter-process switch** causes process idle time: to load + restore PCB
- Idle time generally to be avoided, or minimized
- To **reduce idle time** caused by storing PCB info onto disk and then re-loading PCB from disk . . .
- Leave all PCB info in memory, even of suspended process, so that multiple disk accesses can be skipped during context switch
- Conflict: **less free physical memory** available for other active processes
- Typical for an OS: ***conflicting goals***

CPU Switch *Process to Process*

- Scheduler's **switch** from one process to another via interrupt is similar to function call + return
- Except during *process switch* the changing entities are **unrelated, different processes** being interrupted
- While in the function call + return mechanism all actions are part of the same program
- But after some “distraction”
 - Once the calling function experiences the callee's return . . .
 - Once the interrupted process receives a new time slice . . .
- Execution continues where it had left off before the change in execution flow: **transparently!**

Threads & Scheduling

Threads

- Process often viewed as a **single thread** of execution
- For Single Thread, the **thread notion** is not meaningful!
- Consider executing **multiple threads**, using multiple program counters (pc), **one pc per thread, all progressing as part of one single program**:
 - **Multiple code locations** can sometimes execute simultaneously on MP target; or concurrently on UP target; even that can speed up execution! **But how?**
 - By reducing “wait” conditions that otherwise would have happened, costing real idle time for process; not for CPU
 - **Multiple control sequences** can execute → **threads**
 - And still just one single, original program is executing!
- Can any and every part of process become a thread?

Threads

- Process often viewed as a **single thread** of execution
- For Single Thread, the **thread notion** is not meaningful!
- Consider executing **multiple threads**, using multiple program counters (pc), **one pc per thread, all progressing as part of one single program:**
 - **Multiple code locations** can sometimes execute simultaneously on MP target; or concurrently on UP target; even that can speed up execution! **But how?**
 - By reducing “wait” conditions that otherwise would have happened, costing real idle time for process; not for CPU
 - **Multiple control sequences** can execute → **threads**
 - And still just one single, original program is executing!
- Can any and every part of process become a thread?
- **Only logically independent sections can! Detail later...**

Linux Task

Virtual Memory data structure of Linux:

```
// mm_struct describes virtual memory of a task or process
struct mm_struct {
    int      count;
    pgd_t    * pgd;
    unsigned long context;
    unsigned long start_code, end_code, start_data, end_data;
    unsigned long start_brk, brk, start_stack, start_mmap;
    unsigned long arg_start, arg_end, env_start, env_end;
    unsigned long rss, total_vm, locked_vm;
    unsigned long def_flags;
    struct    vm_area_struct * mmap;
    struct    vm_area_struct * mmap_avl;
    struct    semaphore mmap_sem;
};
```

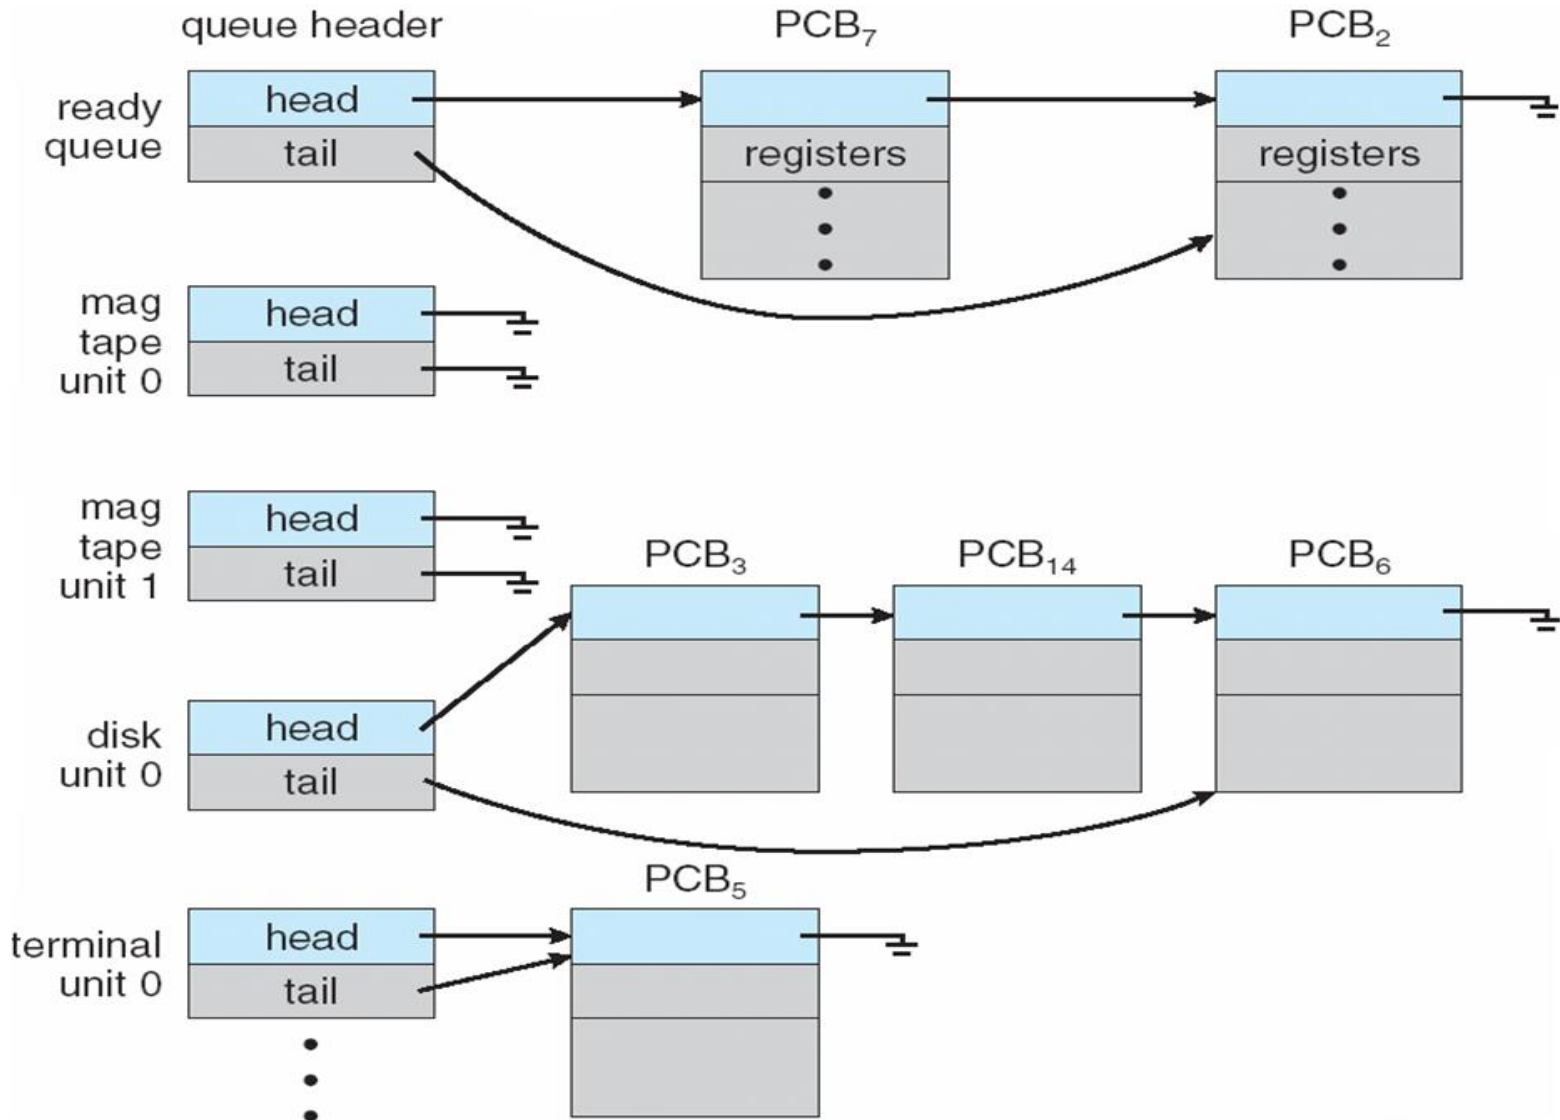
Process Scheduling

- To maximize CPU use, **process switch** must be **fast!**
- Process switch AKA **context switch**
- What is a fast context switch?
 - 1/10 of a second?
 - 1/100 of a second?
 - 1/1000 of a second? (AKA **ms**)
 - 1/10,000 of a second?
- In one single **ms** some million instructions (absent memory-access and file IO) can be executed
- So a context-switch time **as long as a ~ms** seems long, seems to be high overhead
- Typical times: ~100 ns to few μ -s

Process Scheduling

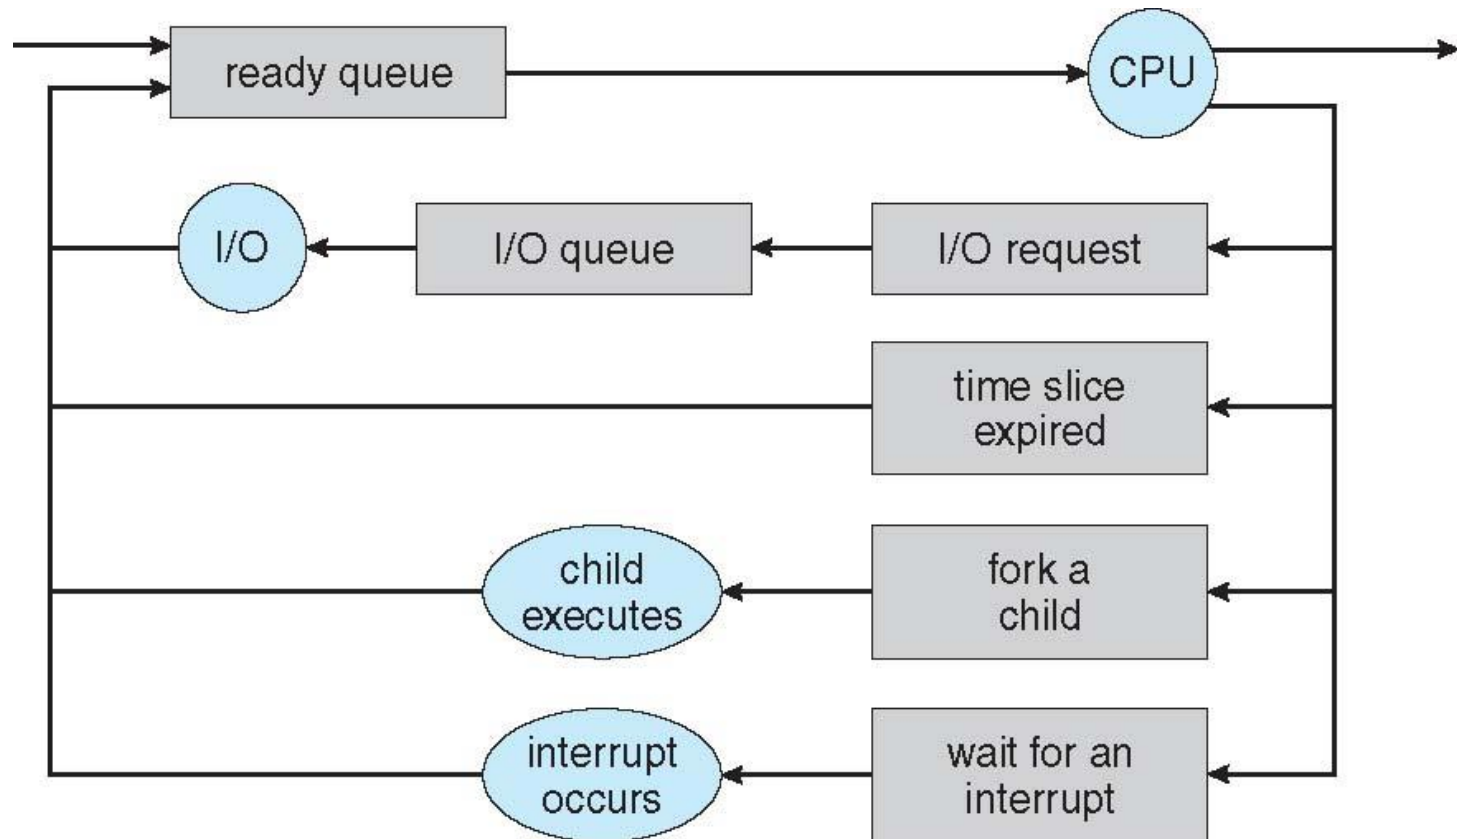
- **Process scheduler** (part of OS) selects via interrupt one among several available processes as next candidate for execution –for up to a full time slice; or less if being interrupted again
- OS maintains **scheduling queues** of processes
 - **Job queue** – all processes not yet completed in system
 - **Ready queue** – queue of all processes in main memory, ready and waiting to execute; just need a CPU time slice
 - **Device queues** – processes **waiting** for device, e.g. for I/O; AKA **wait queues**
 - Other queues as needed
 - Processes generally **migrate** among various queues, and may change priorities, during migration and even while waiting!

Process Queue Scheduling



Process Scheduling

Queuing diagram representing: queues, resources, flows

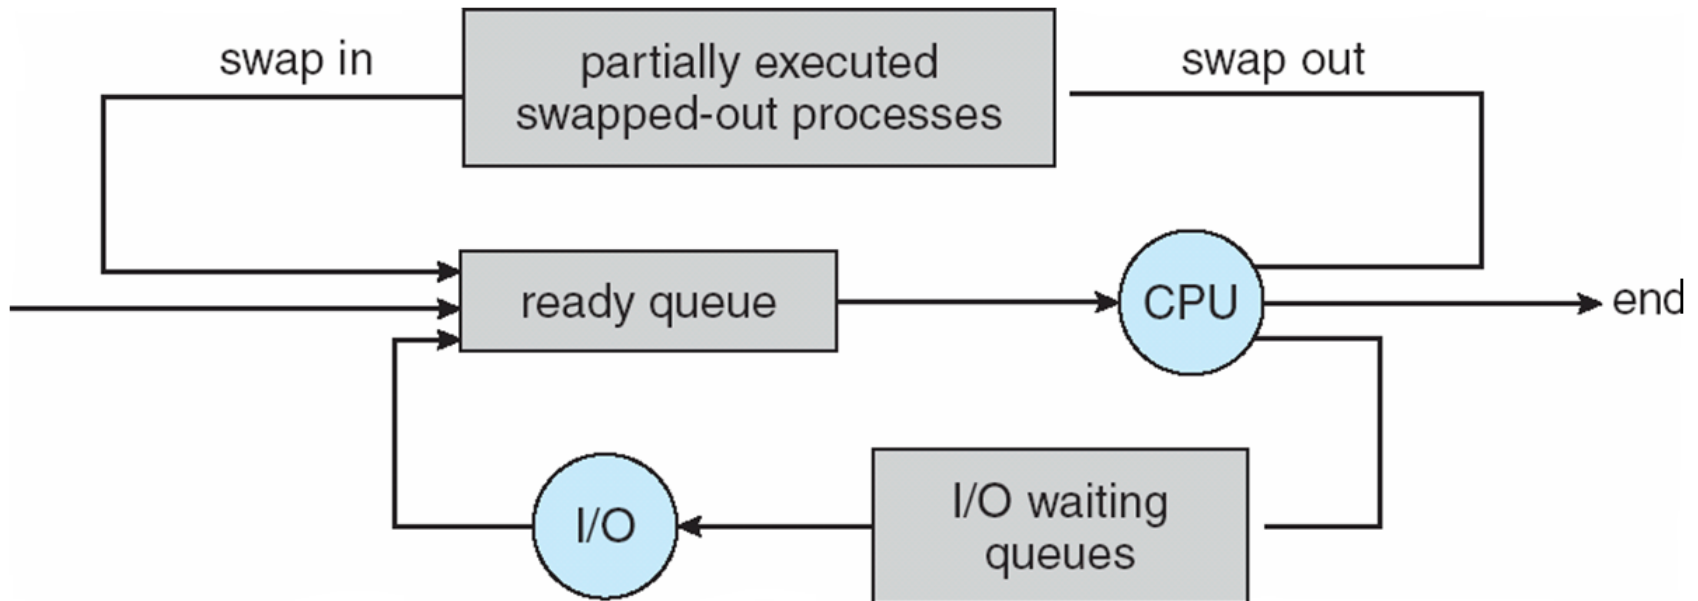


Schedulers

- **1. Short-term scheduler** (or **CPU scheduler**) – selects via interrupt which process to execute next and to grant the CPU to
 - Sometimes this is the only scheduler in a system
 - Short-term scheduler is invoked frequently: ⇒ **fast**
- **2. Long-term scheduler** (or **job scheduler**) – selects which processes, and how many, should be brought into **ready** queue
 - Long-term scheduler invoked rarely: ⇒ **slow** OK
 - Long-term scheduler controls **degree of multiprogramming**
- Processes can be characterized as:
 - **I/O-bound** process – spends more actual time doing I/O than computations: many short CPU bursts
 - **CPU-bound** process – spends more time doing computations; typical for long CPU bursts
- Long-term scheduler strives for good *process mix*

Schedulers

- **3. Medium-term scheduler** optional; can be added if degree of multi programming needs to decrease
- Removes process from memory, stores on disk, swaps back in from disk to continue execution if situation warrants: **swapping**



Multitasking in Mobile Systems

- Some mobile systems –early version of **iOS**– allow **only one process** to run (on Apple); others stay suspended, i.e. not ready!
- Due to screen real estate, user interface poses hard limits on mobile platform: iOS only provides for
 - Single **foreground** process– controlled via user interface
 - Multiple **background** processes– in memory, running, but not being displayed, and with limits
 - **Limits** include single, short task, receiving notification of events, specific long-running tasks like audio playback
- **Android** runs foreground and background

Context Switch

- When CPU switches to another process, OS must **save state** of old process and load the **saved state** of new process: AKA **context switch**
- **Context** of any process is held in its respective PCB
- Context-switch time is overhead: system does no useful user work while switching (is OS bureaucracy 😊)
 - Except enforces **fairness** among users
 - Enables higher **throughput**
 - Complex PCB causes slow context switch! Hence: Simplify!
- Overhead time, AKA **wait time**, depends on HW support:
 - Some HW provides **multiple register sets** → multiple contexts loaded simultaneously with live data; but only **one running**
 - e.g. Intel **hyperthreaded** MPs
 - OS can leave data of interrupted process in memory; inactive, but quick to save (~0 cycles) and to resume (~0 cycles)

Context Switch

Running process will eventually be interrupted!



Processes

Operations on Processes

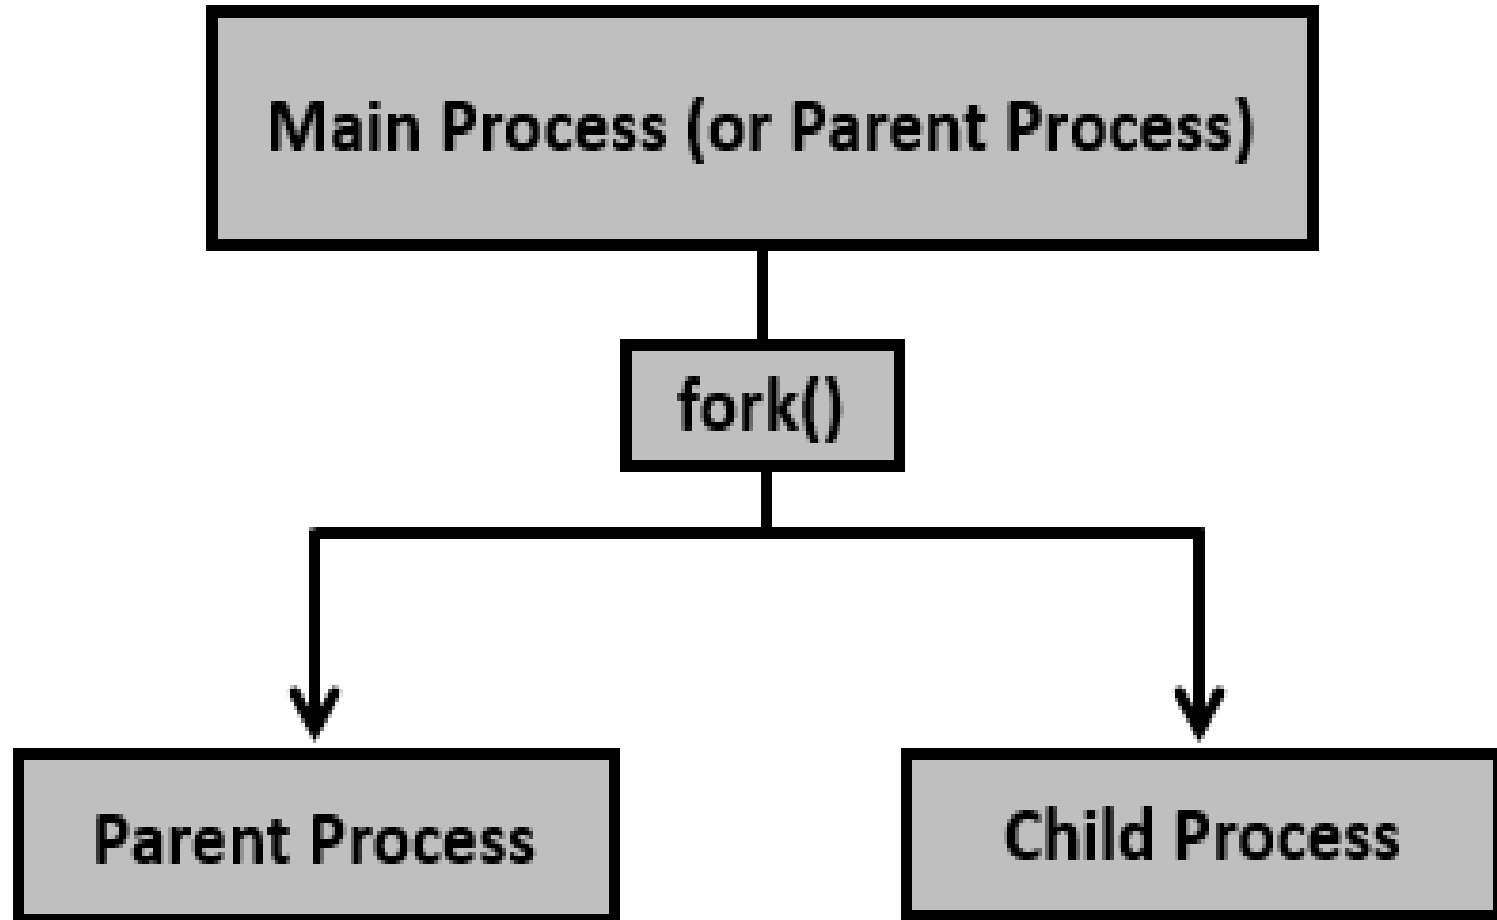
- OS provides mechanism for:
 1. Process **creation**
 2. Process **termination** for various normal and abnormal reasons
 3. Programmed process termination via **exit()** call
 4. Process **status** to be returned to place of call
 5. Also needed: **wait()** to ensure clean end-of-life of child
- Must happen **concurrently** with other processes on the same system, using this **single CPU**
- Effectively competing (**time sharing**) for key resources: services + CPU time, **the** critical resource

Process Creation

- **Parent** process creates **child** processes via **fork()**, which, in turn becomes a new, different process via **exec()**, forming possibly a **tree** of processes
- Generally, process identified and managed via **pid**, short for: **p**rocess **i**dentifier; all **pids** are unique!
- Diverse resource sharing options; OS-dependent, e.g.:
 - Parent and children can share **all** resources
 - Children at times share only **subset** of parents' resources
 - Or parent and child share **no** resources
 - See also: **Capability-based OS model**, [5]
- Execution order
 - Parent and children execute concurrently
 - Parent waits until children terminate; else **zombie** processes

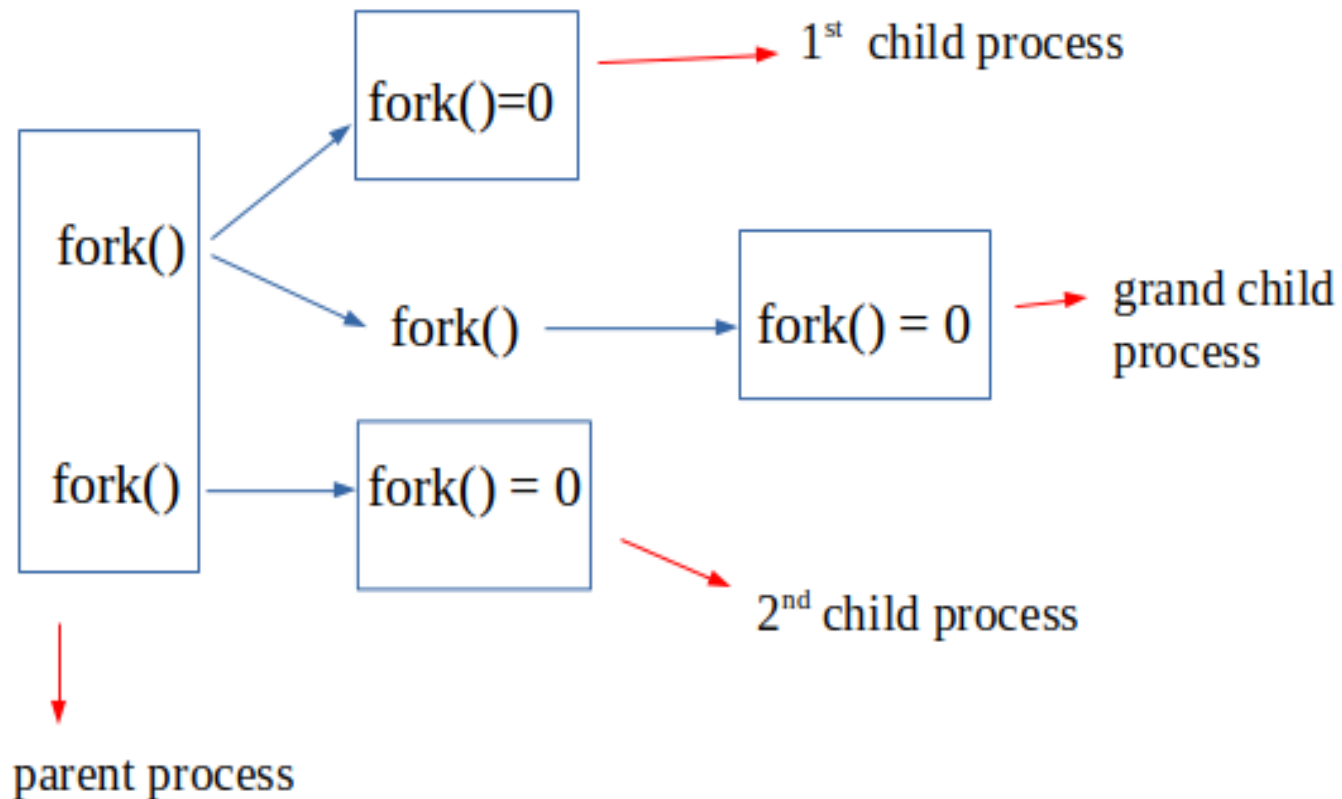
Operations on Processes

Process Creation via **fork()** on Linux & Unix



Operations on Processes

One Program Spawning Multiple Processes via `fork()`



fork () Processes

- **fork()** creates a new process by duplicating the calling process; new process is referred to as the *child process*
- The one originating the fork() is referred to as the *parent* process
- Since parent + child object codes are identical and not modifiable, no need to duplicate!
- At **fork()** time, both memory spaces have identical content –for some time:
- Child process is –until **exec()** call– an exact code duplicate of parent process, except for:


fork () Processes

1. Child has its own **unique process identifier (PID)**, and this PID does not match the ID of any existing process group (**setpgid(2)**) or session
2. The **child's parent's process ID** is simply the parent process ID (PPID)
3. The child does **not** inherit its parent's memory locks (**mlock(2)**, **mlockall(2)**)
4. Process resource use (**getrusage(2)**) and CPU time counters (**times(2)**) are set to **zero in child**
5. The child's set of pending signals is initially empty (**sigpending(2)**)

fork () Processes


6. Child does **not** inherit **semaphore adjustments** from its parent; see [6] on Linux; not covered here
7. Child does **not** inherit process-associated **record locks** from its parent, but does inherit open file description locks from parent
8. Child does **not** inherit **timers** from parent
9. Child does **not** inherit outstanding **asynchronous I/O** operations from its parent, nor asynchronous I/O contexts from parent

A Tree of Processes in Linux



Process Creation Detail

- **Address space**
 - Child is duplicate of parent (modulo exceptions listed above)
 - Child can load new, different code into it's space via **exec()**
- **UNIX examples**
 - **fork()** system call creates new process (a duplicate)
 - **exec()** system call is used after **fork()** call to replace process' memory space with a **new program space**
 - Also **execve()** common on Linux



Process Termination

- Process executes last instruction, then asks OS to terminate it using `exit()` system call
 - Return status from child to parent is passed via `wait()`
 - Process resources are de-allocated by OS
- Parent may also terminate child processes prematurely, using `abort()` system call; abnormal!
- Plausible **reasons for early abort**:
 - Child has **exceeded** allocated resources
 - The task assigned to child is **no longer required**
 - Parent is exiting for external reasons, and the OS **prohibits** a child to continue if parent is terminated

Process Termination

- Some Operating Systems **don't allow child to exist if parent terminates**: consequently, when process terminates, all its children must terminate as well
 - AKA **Cascading** termination
 - All children, grandchildren, etc. are then terminated
 - Termination initiated by OS
- **Parent process may chose to wait** for termination of a child process by using the `wait()` system call

```
pid = wait( & status );
```

Process Termination

- A **zombie** is in existence, when a parent process does not use **wait()** after a child terminates to read that child's exit status; i.e. there is no record that the child ever did die!
- Zombie process **exists via an entry in the OS process table**; i.e. zombie is a process in terminated state, but is still “accounted for”
- An **orphan** is a child process whose original parent process terminates before that child process itself terminates
- Detail:

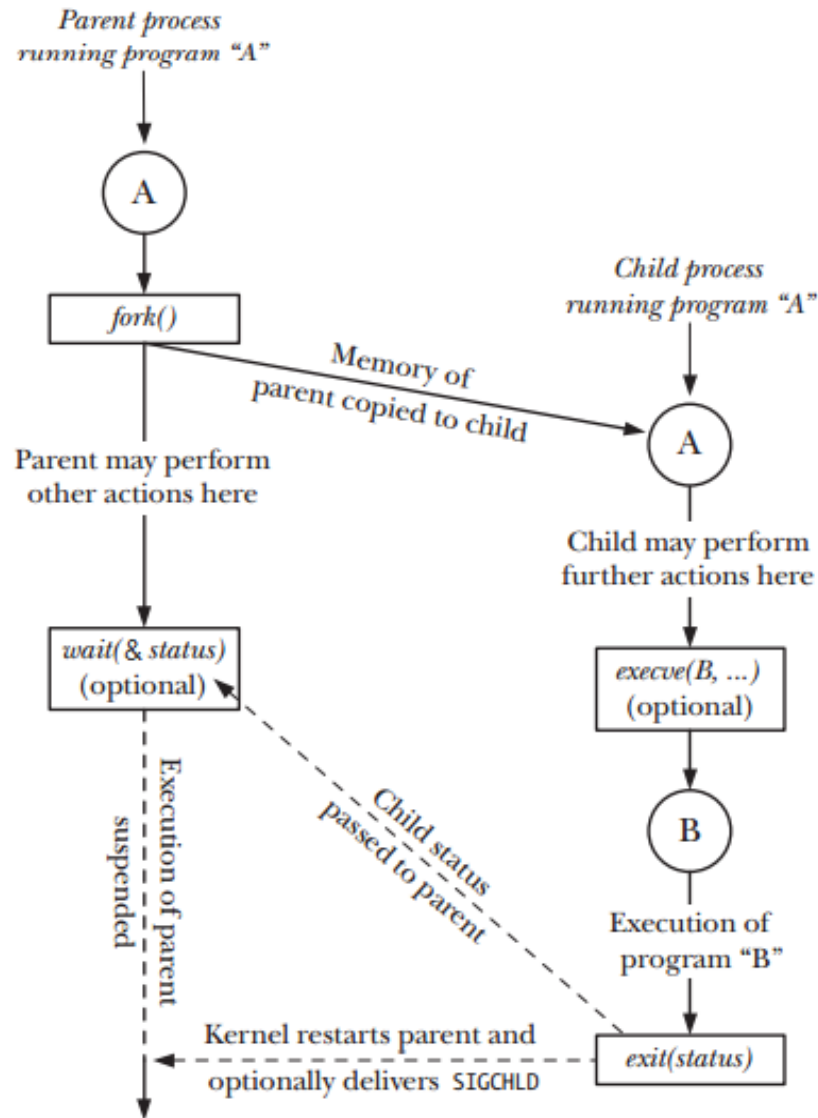
Process Termination


- Under Unix, a **zombie** process or a defunct process is a process that has completed execution via **exit()** but still has an entry in the process table
- A zombie is a process in the **Terminated** state
- This occurs, whenever the entry in the process table is needed to allow the parent to read its child's exit status
- This is generally an **orderly way of executing!**
- Once the exit status is read via **wait()**, the zombie's entry is removed from the process table
- Then the zombie is said to be **"reaped"**

Process Termination

- A child process first **becomes a zombie** before being removed from the resource table
- Generally, zombies are immediately waited on by their parent and then **reaped** by the system
- Processes that stay zombies for a long time are generally an error and cause a minor **resource leak**
- But the only true resource they consume is the one entry in the process table

Process Termination





Unix

- Contemporary **process model** was popularized under Unix
- Key contributors: Dennis **Richie** and Ken **Thompson**
- Here hard ☺ at work on a PDP11 in 1972
- Refining Unix OS
- **Process concept** had already been developed way before Unix

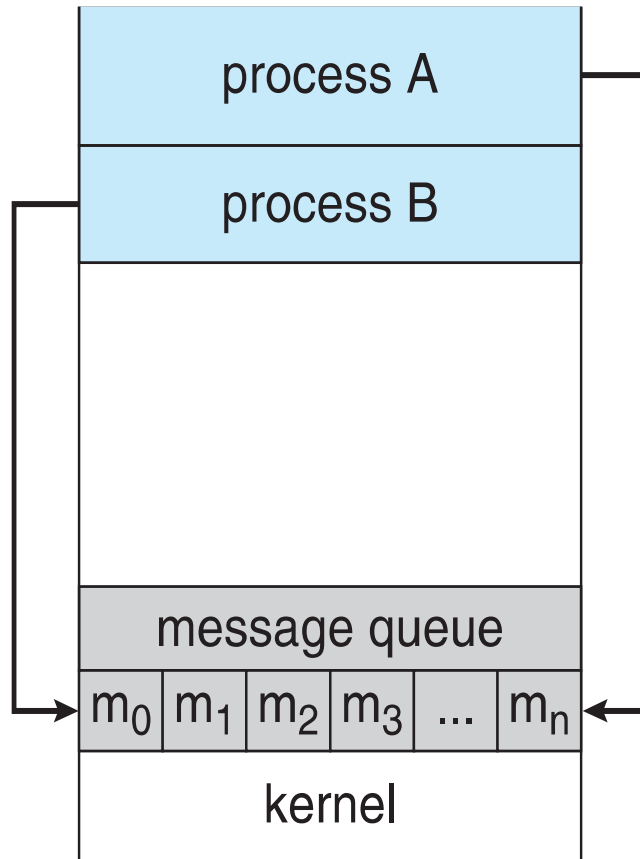
IPC via Shared Memory

IPC Interprocess Communication

- Processes may be *independent* or *cooperating*
- **Cooperating** ones can affect each other
- Reasons for cooperating processes:
 - Information **sharing**, e.g. File sharing
 - Computation speedup
 - Speedup even on UP; if otherwise parent would have to wait for some event to continue computing
- Cooperating processes need so called **interprocess communication, IPC**
- Two communication models of IPC
 - **Message Passing**
 - **Shared Memory**


Communication Models

a) Message Passing



(a)

b) Shared Memory



(b)

Cooperating Processes

- ***Independent*** process cannot affect, or be affected by, execution of another process . . .
- Other than having to **share a key resource**: CPU
- ***Cooperating*** process can affect, or be affected by, the execution of another process, the one it cooperates with

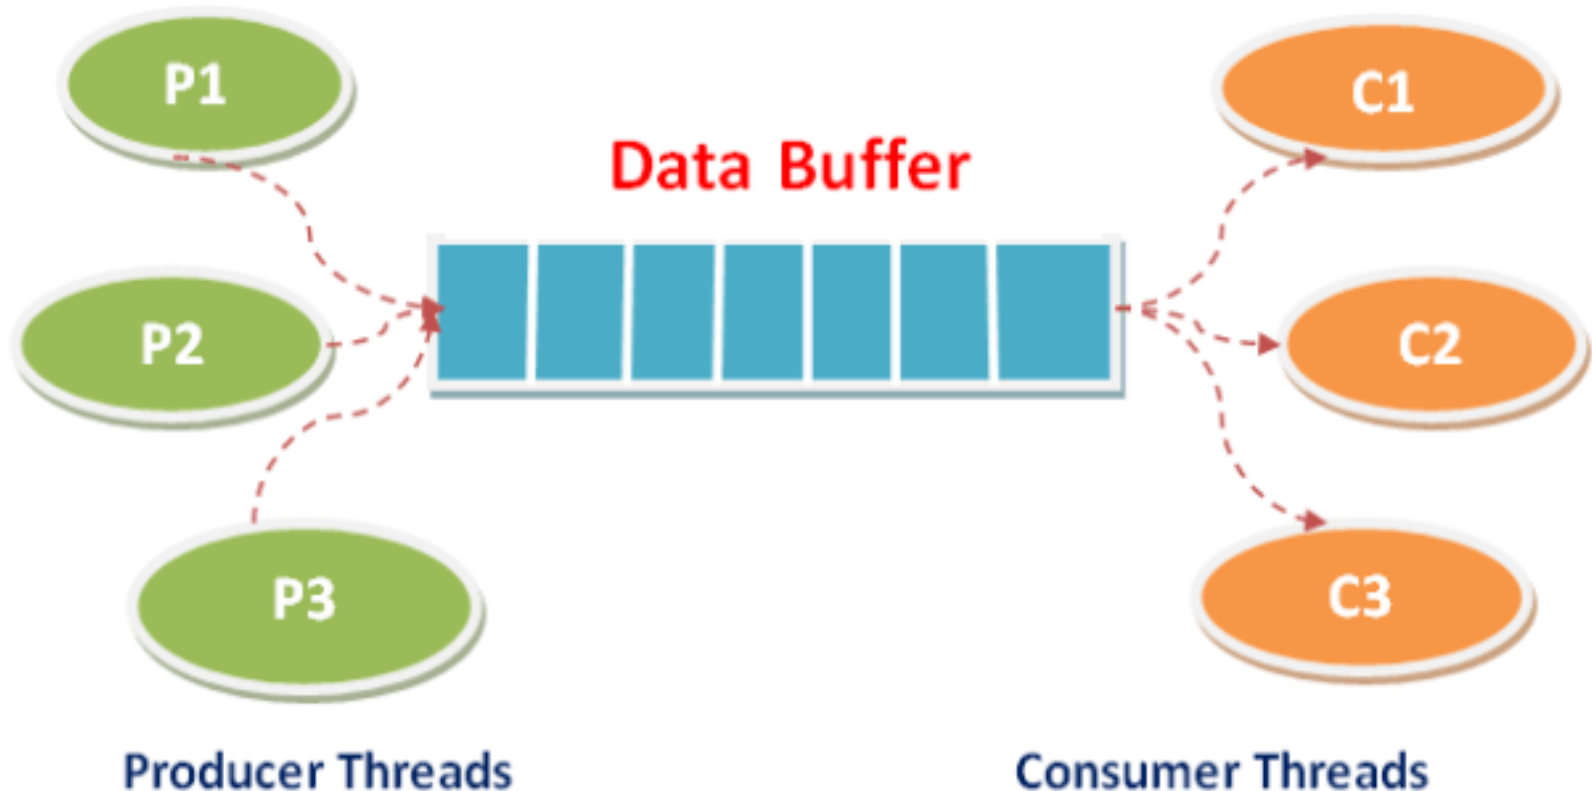
Producer Consumer Problem

- Paradigm for cooperating processes, AKA the so called *producer consumer problem*:
- **Producer** process **generates** information, placed into non-full buffer
 - Note: Producer must **wait()** when buffer is *already full!*
- **Consumer** process **absorbs** such information from a non-empty buffer
 - Note: Consumer must **wait()** when buffer is *still empty!*
- **Unbounded Buffer** places no hard, practical limit on the buffer size; hard to implement ☺ needs infinite space
- **Bounded Buffer** "knows" there is a fixed buffer size; can cause wait() for either producer or consumer

Producer Consumer Problem

Paradigm for Sharing Data Cooperatively

Producer Consumer Problem



The Bounded Buffer Problem

Bounded Buffer, Shared-Memory

- Bounded Buffer paradigm involves processes, and:
 - A *shared data structure*, AKA the **buffer[]**
 - Bounded buffer problem employs **producer** and **consumer** actions. Producer and consumer . . .
 - . . . need data structures: **in** and **out**, indexing next element to be produced, and next element to be consumed in **buffer[]**
 - **Producer** places information into shared **buffer[in]**
 - **Consumer** removes information from shared **buffer[out]**
- Consumer has to **wait**, when buffer is **still empty**
- Producer has to **wait**, when buffer is **already full**
- **In** and **out** are checked and “synchronized” (updated) modulo **buffer size**

Bounded Buffer, Shared-Memory


- Shared data:

```
#define BUFFER_SIZE 10          // shared size e.g. 10
typedef struct {                // type of buffer elem.
    . . .                      // real data here!
} item_tp;                     // arbitrarily complex
item_tp buffer[ BUFFER_SIZE ]; // shared
int  in = 0;                   // should be subrange; shared
int  out = 0;                   // also an "index"; shared
```

- Solution is correct, **but** has a **limited, fixed number** of `BUFFER_SIZE - 1` elements; with loss of 1 index!
- Must reserve one combination –one lost index value– for condition: **Buffer is empty**! Index of the “wasted” slot doesn’t have to be 0! Generally isn’t, only initially


Bounded Buffer Producer

```
item_tp next_produced;           // local to producer
while( true ) {                  // do forever! w. interrupt
    // generate next_produced, place in buffer[ in ]
    next_produced = ...          // Real work here: produce!
    // now check: space in buffer[]?
    while( ( ( in + 1 ) % BUFFER_SIZE ) == out ) {
        ; // wait; full buffer
    } //end while
    buffer[ in ] = next_produced;
    // only producer sets: in
    in = ( in + 1 ) % BUFFER_SIZE;
} //end while
```




Bounded Buffer Consumer

```
item_tp next_consumed;           // local to consumer
while( true ){                   // do forever
    while( in == out ) {        // empty buffer[]?
        ; // wait               // nothing to consume
    } //end while
    // next element available at buffer[ out ]
    next_consumed = buffer[ out ];
    // only consumer sets: out
    out = ( out + 1 ) % BUFFER_SIZE
    // consume next_consumed
    // Real work here: consume!
    . . . = next_consumed
} //end while
```



IPC Shared Memory

- An area of memory **shared** among processes that need to communicate, **occasionally!**
- Communication under **control of user processes**, not under OS control!
- Major goal: Provide a mechanism to allow user processes to **synchronize their actions** when accessing shared memory



Interprocess Communication (IPC) via Message Passing

IPC Message Passing

- IPC is a mechanism for **processes to communicate** and to synchronize their actions, when necessary! Done by sending a message; use a **message system**!
- Generally they **progress independently**! Yet at specific times, they must **cooperate** and **coordinate**!
- Message system means: Processes communicate with each other **without shared objects**
- **IPC** facility provides two message operations:
 - `send(message)`
 - `receive(message)`
- **Message size** can be fixed or variable

IPC Message Passing

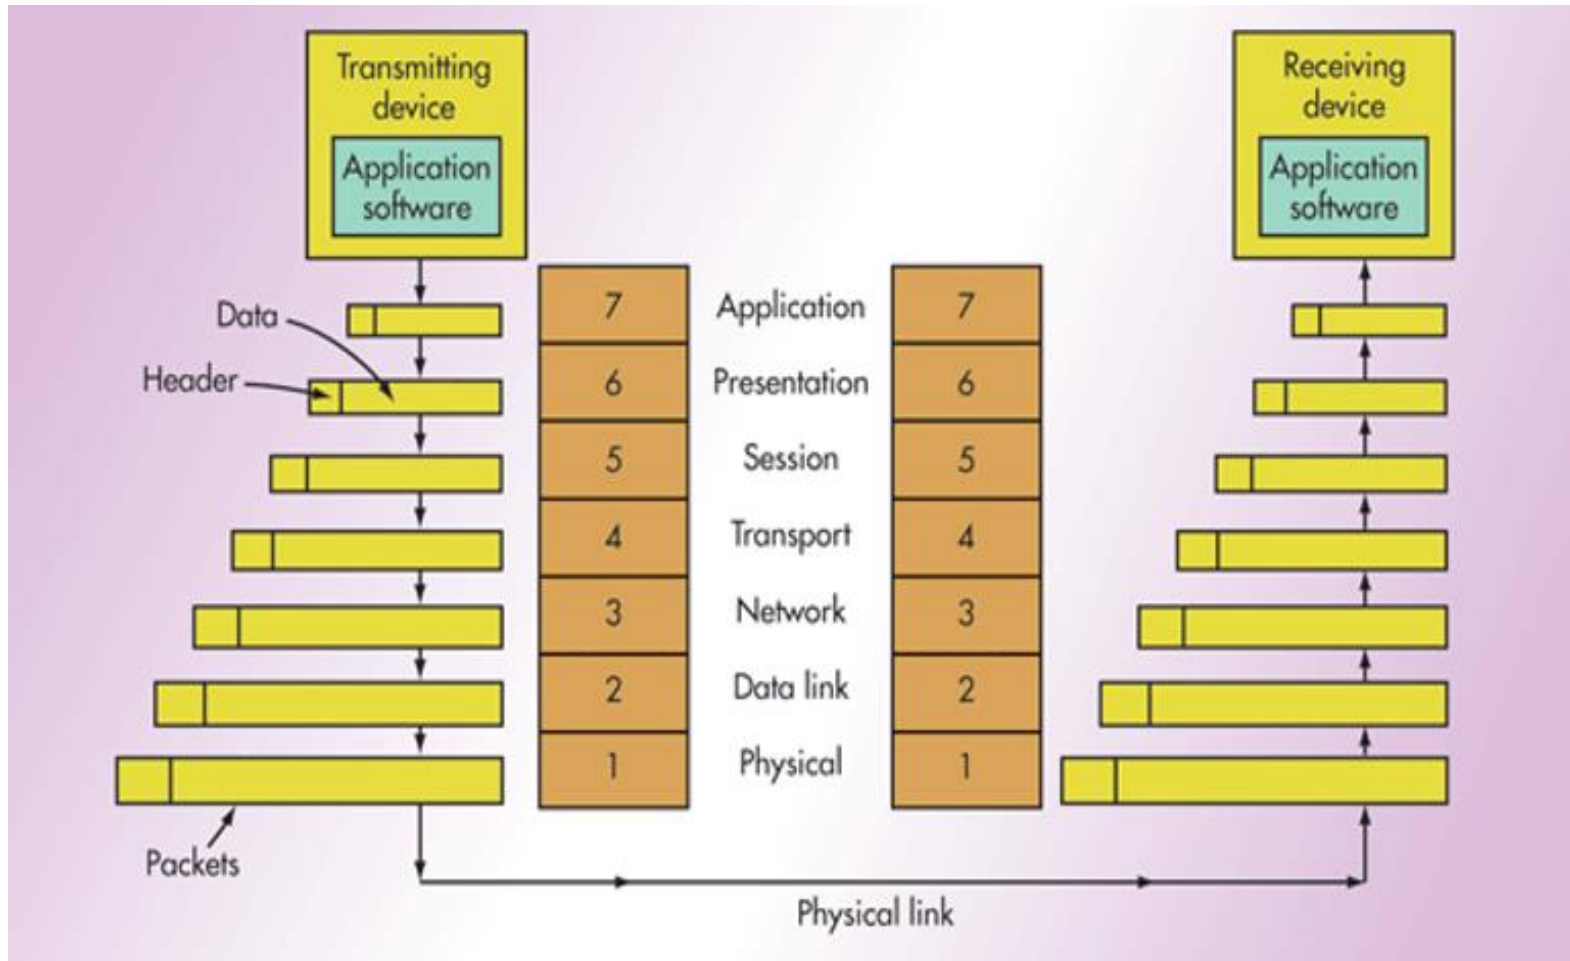
- For processes P and Q to communicate, they need to:
 - Establish a *communication link* between them
 - Exchange messages via `send()` & `receive()`
- Implementation issues and OS design parameters:
 - How are **links established**?
 - Can a link be associated with more than just two processes?
 - **How many links** can there be, between any pair of communicating processes?
 - What is the **capacity** of a link?
 - Is the message size **fixed** or **variable**?
 - Is a link **unidirectional** or **bi-directional**?

IPC Message Passing

- **Communication link** is a **channel** connecting two or more communicating devices
- May be an **physical link** or a **logical link** that uses some other physical connection
- Implementation of communication link:
 - Physical:
 - Special HW bus
 - Network
 - Others
 - Logical:
 - Direct or indirect communication
 - Synchronous or asynchronous
 - Automatic or explicit buffering
- Reminders **TCP/IP**: Transmission Control Protocol, internet protocol. **IPC**: Inter-Process Communication

IPC Message Passing

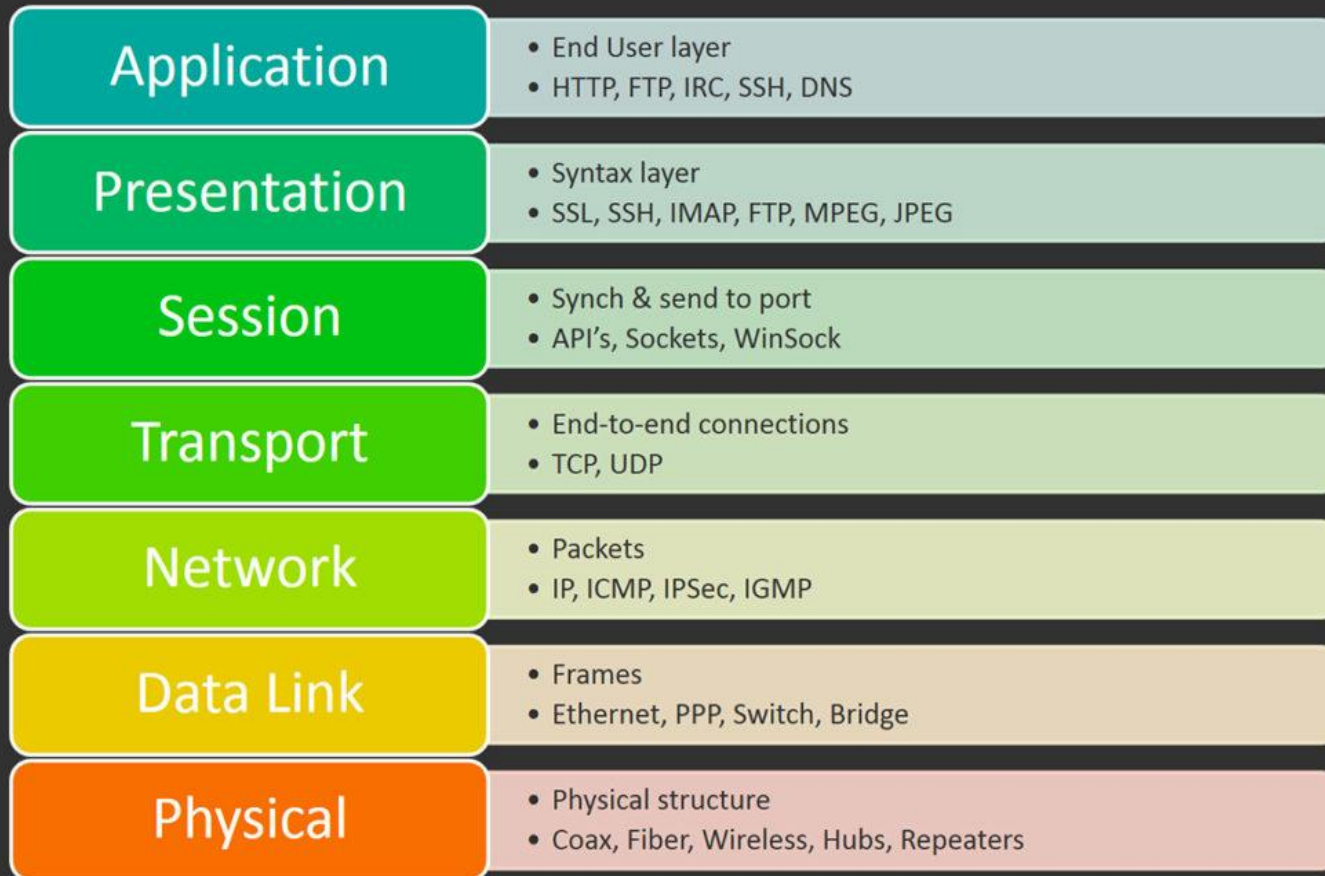
Open Systems Interconnect (OSI): the 7 Layers



IPC Message Passing

Different Image, Same 7 Layer OSI Model

7 Layers of the OSI Model



Direct & Indirect Communication

Direct Communication



Direct Communication

- Processes must **name each other explicitly** in direct communication:
 - `send(P, message)` – send message **to process P**
 - `receive(Q, message)` – receive message **from process Q**
- **Properties of communication link**
 - Links are established automatically
 - A link is associated with exactly **one pair of communicating processes**
 - Between each pair there exists exactly one link
 - Link may be **unidirectional**, but is **usually bi-directional**

Indirect Communication

- Indirect communication involves a **mailbox**
- **Mailbox** being an external device, NOT part of either process!
- But all involved processes can access such a mailbox
- Web on **indirect communication**:
<http://people.scs.carleton.ca/~maheshwa/courses/300/14/node12.html>

Indirect Communication

- Messages are directed to, and received from **mailboxes** (also referred to as **ports**)
 - Each **mailbox** has its own unique **id**
 - Processes can communicate indirectly only if they share such a mailbox
- Properties of **indirect communication** link
 1. Link established only if processes **share a common mailbox**
 2. A link may be associated with more than two processes
 3. Each pair of processes may share several communication links
 4. Link may be unidirectional or bi-directional

Indirect Communication



Indirect Communication

- **Operations**
 - **create** a new mailbox; or port
 - **send** and **receive** messages through mailbox
 - **destroy** a mailbox
- **Primitives are defined as send() receive():**

Note A is a mailbox, not a process ID as on p. 70!

send(*A, message*) send message to **mailbox A**

receive(*A, message*) receive message from **mailbox A**

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Which of the two will receive (AKA read) the message?
- Solution Options
 1. For example, allow a link to be associated with at most two processes; the problem disappears
 2. Allow only one process at a time to execute a receive operation; requires mutex protocol
 3. Allow the system to select arbitrarily the receiver, and then the sender is notified who the ultimate receiver was

Indirect Communication

One thing a Sagittarius hates is indirect communication. They wish more people would just tell it like it is, similar to them.

#SagFACT 🍷 ✓ 100

Synchronization

Synchronization

- **Message passing** may be blocking or non-blocking
- **Blocking** is considered: **synchronous**
 - Blocking **send**: Sender is blocked until the message is received; “Receiver: I can’t read your message now; just wait!”
 - Blocking **receive**: Receiver is blocked until a message is available; “Sender: I have nothing to send yet; just wait!”
- **Non-blocking** is considered: **asynchronous**
 - Non-blocking send: Sender simply sends the message and continues
 - Non-blocking receive: Receiver accepts message
 - A valid message is received, or
 - If no message: continue and check back later, periodically

Synchronization



Synchronization

- **Synchronous serial communication** describes a serial communication protocol in which data are sent at a **constant rate**
- Synchronous communication requires that **clocks** in transmitting and receiving devices **be synchronized** – running at the same rate – so the receiver can sample the signal at the **same time intervals** used by transmitter
- **No start or stop bits** are required. For this reason synchronous communication permits more information to be passed over a circuit per unit time than asynchronous serial communication
- Over time transmitting and receiving clocks tend to **drift apart**, requiring *resynchronization*

Remote Procedure Call (RPC)


Remote Procedure Call

- Remote Procedure Call (RPC) is a protocol that one program, executing **on one computer**, can use to request a service from another program located **on another computer** on some network
- Without having to understand the network's details
- Procedure is AKA **function call** or subroutine call
- **RPC** thus executes the client-server model

Remote Procedure Call

- **RPC** abstracts procedure calls between processes on networked systems
- **Stub** – is SW that converts parameters passed between client and server during RPC
- Client-side stub locates server and **marshals** params
- Server-side stub receives message, unpacks the marshaled parameters, and performs the procedure on the server
- **Marshaling**: Process of transforming the memory representation of an object to another data format suitable for storage or transmission

Remote Procedure Call



Remote Procedure Call

- On Windows, stub codes compile from specification written in **Microsoft Interface Definition Language (MIDL)**
- Data representation handled via **External Data Representation (XDL)** format to account for different architectures
 - **Big-endian** and **little-endian**
- Remote communication has more failure scenarios than local
 - Messages can be delivered *exactly once* rather than *at most once*
- OS typically provides a rendezvous (or **matchmaker**) service to connect client and server


Steps of RPC

- Remote procedure call: **Interprocess communication technique** used for client-server based applications
- Client has a request message that the RPC translates and sends to the server
- This request may be a procedure or a function call to a remote server
- When server receives the request, it sends the required response back to the client
- The client is blocked while the server is processing the call; resumes execution after the server is finished
- The sequence of 6 defined events in remote procedure call is:

Steps of RPC

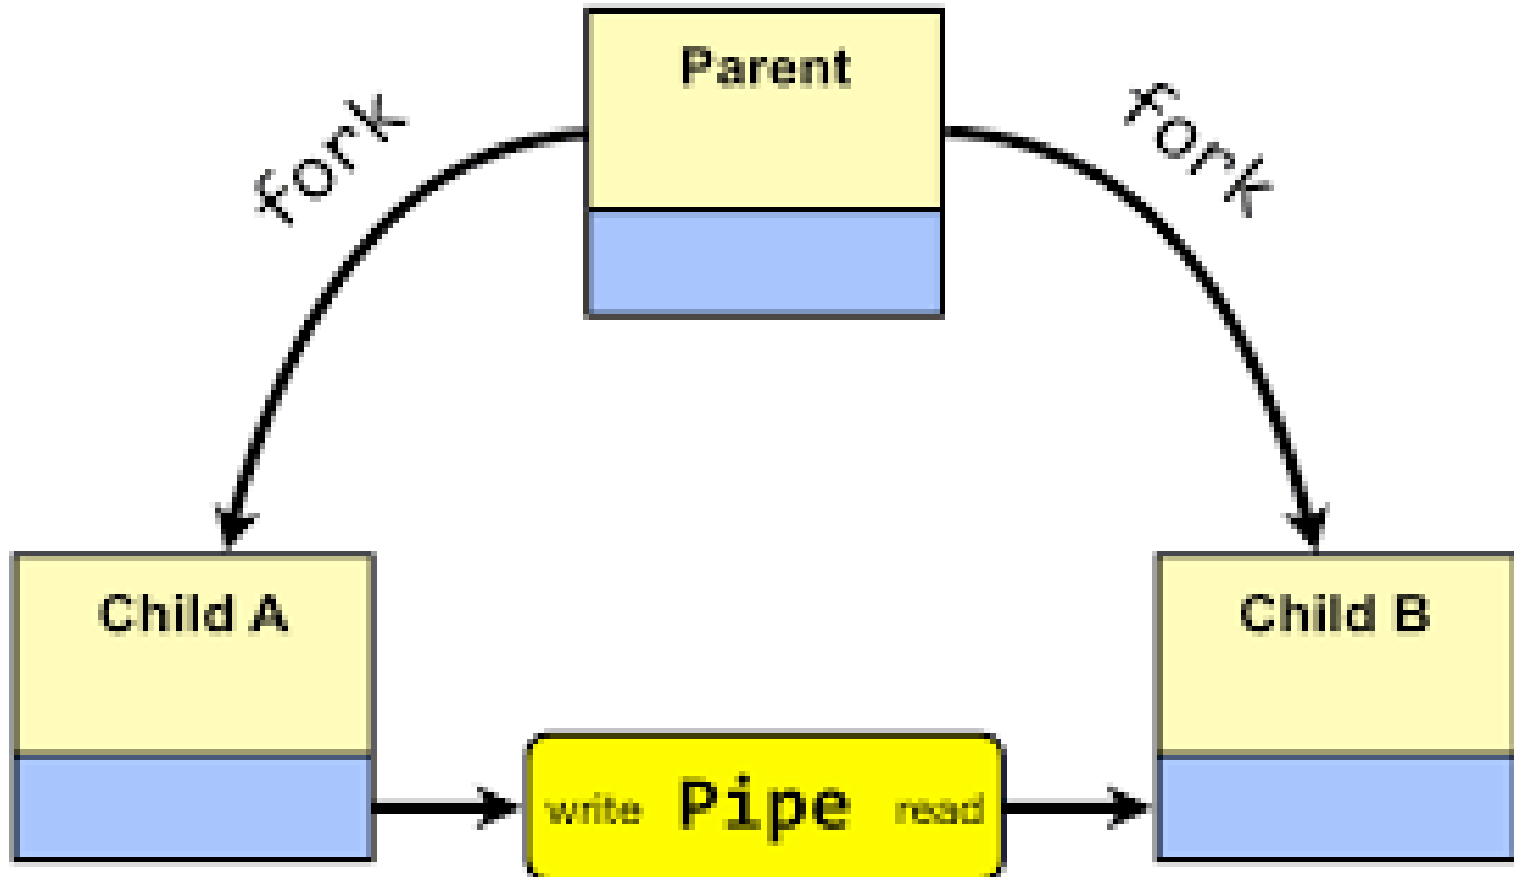
1. Client **stub is called** by client
2. Client **stub** makes a system call to send the message to server, **places parameters** into message
3. Message now is sent from the client to the server by the client OS
4. Message is passed to server stub by the server OS
5. Parameters are removed from the message by server stub
6. Then server procedure is called by the server stub
 - Note the generic stub-to-OS and OS-to-stub mechanism, which allows “detail” to be handled locally; network has no “need to know”

Execution of RPC



Pipes

Pipes




Pipes

- A Unix **Pipe** is a conduit allowing two processes to communicate
- Design options to be decided:
 - Is communication unidirectional or bidirectional?
 - In case of two-way communication: half or full-duplex?
 - Must there exist a relationship (i.e., *parent-child*) between the communicating processes?
 - Can the **pipes be used over a network?**
- **Ordinary pipe** cannot be accessed from outside process that created it
- Parent process creates pipe and uses it to communicate with a child process
- **Named pipe** can be accessed without parent-child relationship

Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end; AKA the **write-end**
- Consumer reads from other end; AKA **read-end**
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes
- Windows calls these **anonymous pipes**



Named Pipes

Named Pipes have added capabilities vs. ordinary pipes:

- **Communication is bidirectional**
- **No parent-child relationship is necessary between the communicating processes**
- **Several processes can use the named pipe for communication**
- **Provided on both UNIX and Windows systems**

Summary

- Process is a **program in execution**, in one of various states of execution!
- In Unix-like environment, OS tracks processes through five-digit ID known as the **pid**, short for **process ID**
- Interprocess communication (IPC) is a programming interfaces to coordinate activities among different processes running concurrently
- Allows program to handle multiple simultaneous user requests
- Remote Procedure Call (RPC) used to request service from a program located in computer on a network
- **Pipe** is a unidirectional data channel used for interprocess communication

Bibliography

1. **Unix process:** <https://www.tutorialspoint.com/unix/unix-processes.htm>
2. **Linux process management:**
<https://www.guru99.com/managing-processes-in-linux.html>
3. **IPC:** https://en.wikipedia.org/wiki/Inter-process_communication
4. **C data structure for Linux VMM:**
<https://www.tldp.org/LDP/tlk/ds/ds.html>
5. **Capability Model:** https://en.wikipedia.org/wiki/Capability-based_security
6. **Semaphore adjustment:** <http://man7.org/linux/man-pages/man2/semop.2.html>